

Something Simpler Wins, But Not Always the Same One: Auditing RLVR Control Groups Across Three Backbone Scales

Extended version: complete per-group statistics, aggregator sensitivity
and a selection-bias audit of the decision rule

Keyi Li Yihao He Quanyi Li Guanghui Liu Xiaoyu Lu
Linying Jiang*

Software College, Northeastern University, Shenyang, China

September 20, 2026

Abstract

Reinforcement learning with verifiable rewards (RLVR) is evaluated against a remarkably uniform control group: the base model, matched supervised fine-tuning, sometimes rejection-sampling SFT. All of these are learning arms whose accuracy grows with the demonstration budget. We pre-registered, before the second task was run, a decision rule comparing GRPO against the *maximum* over a four-arm control group, one arm of which is a deterministic symbolic solver that does not learn. Across two program-annotated financial QA tasks and three backbones spanning $14\times$ in scale, all 17 frozen-test effect sizes are negative, from -1.06 to -16.42 points. The mechanism is a moving baseline: the solver attains the maximum at 0.5B and 1B because it is flat in the demonstration budget, while at 7B it falls 15 points behind and continued SFT takes its place. Ablations, a verifier-hacking audit over 57,517 positive-reward samples and a KL sweep all leave the sign unchanged.

One boundary, found by us and reported here. Config 3 was included to answer “your models are too small.” It cannot. Re-running its smallest-margin cell at ten times the prompt coverage reverses it decisively: GRPO reaches 87.23% and 86.88% against continued SFT’s 80.50% ($\Delta = +6.74, +6.38$; McNemar $p = 0.0013, 0.0029$), which clears this paper’s own Go threshold. At the *registered* budget nothing changes, and a run that fixes the degeneracy we diagnose lands at $+1.06$ with an interval of $[-2.13, +4.26]$ — statistically the same place as the original -1.06 . So the 7B result marks a *budget* boundary, not a capability one, and the quantitative core of this paper is the 11 groups at 0.5B and 1B. Details in Appendix G.

The rest of this extended version adds the material a four-page paper cannot carry, and it is deliberately two-sided. In our favour: the max in the decision rule is *not* responsible for the negative sign—a cross-fitted estimator that selects the control arm on one half of the test set and scores it on the other still yields a negative Δ in 17 of 17 groups (mean -4.18 against a naive -4.44), and the bootstrap selection bias never exceeds 0.47 points; the inference-time sampling arm that is usually proposed as the missing baseline is not deployable on these tasks, and even its oracle upper bound (pass@16) loses to the solver at three of four operating points. Against us: under Holm correction across all 17 groups only **2** are significant, the per-group bootstrap intervals exclude zero in only 5, and **no GRPO run completed even one epoch** over its prompt pool (0.10–0.79 epochs; at 7B roughly 10%). We report all of it, with every number regenerated from the stored per-item outcome vectors by a single script.

*Corresponding author: jiangly@swc.neu.edu.cn

Contents

1	Introduction	3
2	Related Work	4
3	Experimental Setup	6
4	Five Arms and the Decision Rule	7
5	Results	8
5.1	Every frozen-test effect size is negative	8
5.2	Why: the scarce regime is the non-learner’s regime	8
5.3	The gain is real; the conclusion is not	9
5.4	Dev selection bias, observed three times	9
6	Scale: the winning baseline changes identity	10
7	Ruling Out the Obvious Objections	11
7.1	“The solver is just retrieving gold answers”	11
7.2	“Any non-learning baseline would look that strong”	11
7.3	“GRPO’s score is inflated by verifier hacking”	12
7.4	“You did not tune GRPO properly”	12
7.5	One excluded run	13
7.6	Is our GRPO simply under-trained?	13
7.7	“You left out the obvious inference-time baseline”	16
7.8	How the solver was built, and what it cost	16
8	The Decision Quantity, Audited	17
8.1	The pre-registered question: does GRPO clear +3?	17
8.2	Per-group significance, and what it does not show	18
8.3	What the sign pattern is worth	19
8.4	Is the max itself producing the sign?	19
8.5	How much of the result is the solver?	21
8.6	The two arms are not solving the same problems	21
9	Discussion	21
10	Conclusion	24
A	Every frozen-test number	27
B	Complete solver ablation	27
C	Non-learning alternatives: prompting and sampling	28
C.1	Program-of-Thought / PAL prompting	28
C.2	Best-of- n and majority voting	28

D	Verifier-hacking audit	29
E	Training stability across all 18 GRPO runs	30
F	Pre-registration, deviations and reproduction	30
F.1	What was fixed, and when	30
F.2	Artifacts	32
F.3	The inference-engine equivalence gate	33
F.4	Reproducing the numbers	34
F.5	Hardware and software	34
G	Follow-up: where the 7B result ends	35
G.1	Design	35
G.2	What happened	35
G.3	What this does to the paper	36

1 Introduction

Reinforcement learning with verifiable rewards (RLVR) has become the default recipe for improving symbolic reasoning: sample programs, execute them, reward the correct ones, optimize with a policy-gradient method such as GRPO [1]. Reported gains are large and consistent [2].

What those gains are measured against is remarkably uniform. Across the empirical RLVR literature the control group consists of *learning* arms: the un-tuned base model, a matched supervised fine-tuning (SFT) run on the same demonstrations, and—in the more careful papers—rejection-sampling SFT [3, 4]. Every one of these consumes the demonstration budget and improves as it grows. A fourth option is almost never present: a deterministic, non-learning symbolic solver. For program-annotated numerical reasoning benchmarks such as FinQA [5] or TAT-QA [6], where answers are short arithmetic programs over a table, such a solver is not exotic—it is the classical approach the neural pipeline replaced.

This paper asks whether that omission changes conclusions, and finds something more general than expected.

The observation that starts it. With two demonstrations on CodeTAT-QA, warm-start SFT reaches 19.0% dev accuracy and GRPO reaches 72.0%—a 53-point gain that, against a warm-start-only control, is an unambiguous success. A rule-based solver handed the *same* two demonstrations—and using neither—also reaches 72.0%. On the frozen test set the three seeds give $\Delta = -2.13, -1.42, -3.90$ points. The gain over the learning baseline is real; the conclusion drawn from it is not.

Design. We pre-registered—before running anything on the second task—a single decision quantity:

$$\Delta(m) = \text{Acc}(\text{GRPO}|W_m) - \max_{a \in \mathcal{B}} \text{Acc}(a), \quad (1)$$

with $\mathcal{B} = \{W_m, \text{c-SFT}, \text{RS-SFT}, \text{solver}\}$, all arms branching from the same warm start W_m trained on the same m demonstrations. Taking a max over the control group, rather than comparing against one chosen opponent, is the central design decision: it makes the comparison impossible to ease by dropping an arm.

Findings. We evaluate two tasks and three backbones spanning $14\times$ in scale. All 17 frozen-test groups give negative Δ , from -1.06 to -16.42 points (Figure 2). But the interesting structure is in *which* arm wins:

- At **0.5B** and **1B**, the maximum is attained by the deterministic solver in every one of 11 groups. Dose-response curves explain why: the solver is flat in the demonstration budget while every learning arm scales steeply with it, so the scarce regime is where a non-learning baseline is strongest.
- At **7B**, the solver drops 15 points behind and is no longer competitive—exactly as the “RLVR gains emerge with backbone capability” account predicts. Δ stays negative in all six groups anyway, because the winning arm becomes **continued SFT**: the most elementary baseline in the set, and one that appears in essentially every RLVR paper.

So the finding is not “a symbolic solver beats RL”—that is a fact about one scale regime. It is that *at every scale we measured, some simpler method matched or beat GRPO, but which method that is changes with scale*. A control group holding only the base model and one matched SFT run misses the solver at 0.5B; a control group designed around the solver misses continued SFT at 7B. What the evidence demands is breadth of the control group and a max aggregation over it, not the presence of any particular arm.

Contributions. (i) A control-group design for RLVR evaluation—a non-learning arm plus max aggregation—together with the pre-registered decision rule built on it; (ii) an execution over two tasks and three backbones in which all 17 frozen-test effect sizes are negative and the identity of the winning baseline shifts with scale; (iii) three independent observations of dev-selection bias, each a non-negative Δ on dev that turns negative on the frozen test.

Section 8 additionally audits the decision rule against the selection bias it could itself introduce, and finds that it does not account for the result. We report that as an audit of our own instrument rather than as a contribution, alongside the per-group statistics—two of seventeen groups significant under Holm—that the effect sizes actually support.

What we do not claim. We do not claim RLVR is ineffective: our GRPO runs beat every learning baseline at 0.5B, by 35 points at $m=2$. We do not claim the solver generalizes—it is task-specific and its engineering cost is reported in Section 7.8. We do not claim the 7B groups establish a scale effect: at ten times the prompt coverage the smallest-margin 7B cell reverses to $\Delta = +6.74$ (Appendix G), so Config 3 marks a budget boundary rather than a capability one. We claim only that an RLVR conclusion established without a broad control group, a stated budget and a frozen test set is not trustworthy.

2 Related Work

RLVR and its control groups. GRPO [1] and its descendants [2] optimize against an executable checker rather than a learned reward model, which removes reward hacking of the preference-model kind and makes the training signal cheap on program-annotated tasks. Empirical reports pair RLVR against the base policy and a matched SFT run; stronger evaluations add rejection-sampling SFT or iterated self-training [7, 3, 4], which use the same verifier at training time without policy-gradient updates. Every arm in these control groups is a learning arm whose accuracy grows with the demonstration budget. Our contribution is not a new algorithm but the observation that adding one non-learning arm changes the sign of the reported effect.

Recent critical work on R1-Zero-style training has begun to audit the recipe itself rather than its reported margins. Dr. GRPO [8] identifies an optimization bias in GRPO’s length normalization, and SimpleRL-Zoo [9] shows how sharply the outcome of zero-RL training depends on the base model across ten back-

bones. Both findings concern what happens *inside* the RL arm. Our question is orthogonal and complementary: holding the RL arm fixed, what does the control group have to contain before its reported margin means anything? A paper can adopt Dr. GRPO’s unbiased estimator and still report a gain that a non-learning arm erases.

Does RL exceed the base model’s reach? A parallel line asks whether RLVR expands a model’s capability or merely sharpens its sampling distribution, e.g. by comparing $\text{pass}@k$ before and after RL [10] or by auditing where reported gains come from [11]. That question is about the policy; ours is about the comparison. The two can agree in direction while differing in what would refute them: our claim would be refuted by a strong solver that still loses to GRPO, not by a demonstration that RL narrows the output distribution.

Weak baselines as a source of illusory progress. The stance we adopt has an explicit lineage, which the conference version of this paper invoked without citing. In recommendation, Ferrari Dacrema et al. [12] reproduced only 7 of 18 recent neural methods and found 6 of those 7 beaten by properly tuned classical baselines. In deep metric learning, Musgrave et al. [13] found a decade of reported gains largely attributable to experimental protocol. In neural retrieval, [14] made the same argument about comparisons against untuned BM25. Deep RL has its own version of the problem: Henderson et al. [15] showed that reported margins can be smaller than the variation induced by random seeds, and Agarwal et al. [16] showed that point estimates over few runs routinely misrepresent aggregate performance.

We add two elements specific to RLVR. First, the baseline we restore is *non-learning*, so—unlike a tuned BM25 or a tuned matrix factorization—it cannot be strengthened by spending more of the demonstration budget, which is what makes it bind precisely in the scarce regime where RLVR is argued to be valuable. Second, the decision rule is fixed by pre-registration before the second task is run, so the comparison cannot be retrofitted to the result.

Test-set reuse and selection. Our design reserves a frozen test set and evaluates each arm on it exactly once, which is the operational form of a long-standing concern about adaptive analysis. Dwork et al. [17] give the theory of how repeated querying of a holdout destroys its validity; Recht et al. [18] give the empirical counterpart on ImageNet. Our Section 5 reports three instances of the same phenomenon inside a single study, and Section 8 audits our own decision quantity for the selection bias it could itself introduce. For multiple comparisons we use Holm’s sequentially rejective procedure [19], and for the interpretation of non-significant per-group results we follow Card et al. [20] on the power that small NLP test sets actually afford.

Symbolic and program-based numerical reasoning. Program-of-Thought [21] and PAL [22] delegate arithmetic to an interpreter, which is what makes verifiable rewards natural on these tasks. The same property makes a purely deterministic solver feasible: FinQA [5], ConvFinQA [23] and TAT-QA [6] annotate answers with short executable programs over semi-structured tables, and BizBench [24] packages CodeFinQA and CodeTAT-QA in that form.

The observation that policy gradients can be trained against an executable checker on exactly this kind of table-program task is not new, and predates RLVR by most of a decade. Neural Symbolic Machines [25] trained a sequence-to-program model against a Lisp interpreter with REINFORCE under weak supervision; MAPO [26] added a memory buffer of verified trajectories to reduce gradient variance, on WikiTableQuestions [27] and WikiSQL. The control-group problem we describe is visible in that literature too, where deterministic semantic parsers remained competitive baselines. Later, Binder [28] reached state of the art on the same table-QA benchmarks with no training at all, by binding a language model into a symbolic program—the strongest existing evidence that on table-program tasks the non-learning route is a serious

competitor rather than a straw arm. Our solver is a port of classical label-alignment and operator-template machinery, not a novel method, which is exactly why its omission from RLVR control groups is hard to justify.

Pre-registration and frozen evaluation. Pre-registration, held-out test sets touched once, and explicit deviation logs are standard in fields where negative results must be credible. They are rare in ML practice but disproportionately valuable for a paper whose main result is a null: without a decision rule fixed in advance, a reader cannot distinguish “the effect was absent” from “the analysis was run until the effect was absent.” Appendix F states what was frozen when, which comparisons are confirmatory and which are exploratory, and where the full protocol and its logged deviations can be read.

3 Experimental Setup

Tasks and models. We use three configurations built on BizBench [24], chosen so that the dataset family, the model family, the backbone scale, and the strength of the deterministic baseline all vary:

- **Config 1:** CodeTAT-QA $\times$ Qwen2.5-0.5B-Instruct. Context is a JSON table; gold programs index cells as `df["col"]["row"]`. 2846 clean training items, 400 held out as dev (all selection uses a fixed 200-item subset of it), frozen test $n=282$.
- **Config 2:** CodeFinQA $\times$ Llama-3.2-1B-Instruct. Context is narrative financial-report text with embedded tables; gold programs inline the numbers. 4035 clean training items, 400 dev (same 200-item convention), frozen test $n=664$.
- **Config 3:** CodeTAT-QA $\times$ Qwen2.5-7B-Instruct. Identical task, data splits, budgets, and protocol to Config 1 with a $14\times$ larger backbone; this is the scale probe of Section 6.

An item is *clean* if its gold program executes to its gold answer under our verifier: $2846/2856 = 99.65\%$ for CodeTAT-QA and $4035/4668 = 86.4\%$ for CodeFinQA. The ten CodeTAT-QA failures were inspected individually and are dataset defects (a final line not assigned to `answer`, or a comprehension with no bare binary operation), not adaptation artifacts.

Demonstration budgets. Nested subsets are drawn from the clean pool, so a smaller budget is always a prefix of a larger one. Config 1 uses $m \in \{0, 1, 2, 4, 8, 16, 64, 256, 1024, 2446\}$ and Config 2 uses $m \in \{0, 16, 64, 128, 256, 1024, 3635\}$, where the largest value is the full pool. Three independent subset seeds are drawn per configuration; every arm at a given (m, seed) sees exactly the same m demonstrations.

Pre-registration. The decision quantity in Eq. (1), the gate criteria, the budget selection rule, and the verdict vocabulary were frozen before any training or evaluation run on the second task. The protocol also fixed, in advance, how a result contrary to our expectation would be handled: *if the solver is strong and GRPO nevertheless wins, report it as such, and do not weaken the solver to recover the original narrative*. Both outcomes were accepted at freeze time. The solver’s construction was likewise constrained ex ante to a “good-faith maximum effort” clause forbidding template pruning or threshold tuning in response to intermediate results.

Budget selection. Two operating points are chosen per configuration from the warm-start dose curve using the pre-registered bands: m_L is the smallest budget landing in $[5\%, 20\%]$ dev accuracy and m_H the smallest

in $[25\%, 45\%]$. This yields $m_L=2, m_H=64$ for Config 1 and $m_L=128, m_H=256$ for Config 2. When the initial grid left the low band empty, additional budgets were run rather than reinterpreting the band—this is one of the logged deviations.

Leakage control. We assert disjointness of dev, demonstration pool, and test, and verify before each solver run that its retrieval library does not intersect dev. The frozen test set is untouched during training, checkpoint selection, and budget selection; every arm is evaluated on it exactly once, at the checkpoint already selected on dev.

Deviations. Minor protocol adjustments (e.g., adding intermediate budget points when initial grids left required bands empty) were logged alongside their influence direction.

4 Five Arms and the Decision Rule

All arms branch from a common warm start W_m : a LoRA SFT run on the m demonstrations of that budget and seed. This makes the comparison about *what to do next with the same m demonstrations*, not about different amounts of supervision.

The four baseline arms. W_m itself is the starting point. **Continued SFT** keeps imitating the same m demonstrations under the same optimizer settings. **RS-SFT** samples from the current policy on pool questions whose gold programs are *not* revealed, keeps verifier-accepted trajectories, and fine-tunes on them, iterating. **The solver** is a deterministic program: it parses the table, aligns question tokens to row and column labels through a cascade (exact $\rightarrow$ normalized $\rightarrow$ fuzzy $\rightarrow$ year/numeric $\rightarrow$ synonym), classifies the arithmetic intent, and emits a program. Its *scarce* variant—the protocol-compliant one—may retrieve from exactly the m gold programs of that budget, matching the demonstration budget of the learning arms. A *full* variant with access to the whole pool is computed for diagnosis only and never enters $\mathcal{B}$.

The tested arm. GRPO [1] starts from the same W_m and trains on pool questions with the verifier as reward, under a wall-clock budget matched to continued SFT and RS-SFT (Config 1: GRPO 5701–5728 s, c-SFT 5703–5707 s at $m=64$, RS-SFT 4561–5815 s). We note one budget asymmetry in the opposite direction of our conclusion: at $m=2$ in Config 1, continued SFT converged in 2579–3196 s against GRPO’s 5701–5703 s, so GRPO received more compute than the baseline it is compared with.

Implementation and compute. All runs use PyTorch 2.10, transformers 4.57, peft 0.19 and trl 1.8 on a single NVIDIA RTX A6000 (48 GB) per run. All generation uses HuggingFace generate throughout: we tested swapping in vLLM and found it shifts accuracy by 3.00 points—the same magnitude as our Go threshold, and in no systematic direction—so the substitution was rejected (Appendix F.3). Every arm trains a LoRA adapter ($r=16, \alpha=32$, dropout 0) at learning rate 10^{-5} ; GRPO uses 8 generations per prompt, KL coefficient $\beta=0.04$, temperature 1.0, top- p 0.95, and effective batch 64 (per-device 16, accumulation 4), reduced to per-device 4 with accumulation 16 at 7B to fit the card. Each arm is capped at 5700 s of wall clock, enforced at step boundaries, so runs may overshoot by one step. Checkpoints are selected on dev; the test set is evaluated once per group. Section 7 reports the sweep over β , the only hyperparameter we varied after the protocol was fixed.

Decision rule. $\Delta(m)$ follows Eq. (1). Comparison uses integer counts, $k_{\text{GRPO}} - \max_a k_a \geq \lceil 0.03n \rceil$, to avoid floating-point boundary effects at the +3-point Go threshold. Per-item agreement is tested with McNemar and corrected with Holm. Only “GRPO vs. the strongest non-RL arm” is confirmatory; everything else is descriptive.

Where the threshold comes from matters for reading a null result, so we state it rather than assert it. It was fixed between two anchors. The upper one is the published literature. In the RLVR-for-finance papers we could compare against, the gain from adding RL on top of SFT, measured on program-annotated financial QA, runs from -2.56 to $+4.0$ points: Fin-R1 reports $+3.0$ on FinQA and $+4.0$ on ConvFinQA [29], Fino1 $+1.25$ on FinQA [30], and DianJin-R1 -2.56 on FinQA [31]. A bar at five points would therefore have meant “better than anything published on this task family.” We restrict the comparison to program-annotated QA deliberately and note what that excludes: the one published figure we found that reaches five rather than falling below it is Fin-R1’s $+5.0$ on a benchmark that is not program synthesis. Within the task family this paper studies, nothing published clears the bar. The lower anchor is measurement resolution: we estimated single-seed standard error at three to five points, which makes $+3$ the floor of what this design can distinguish from seed noise at all. We set the Go criterion at the lower anchor, $+3$, which is the more permissive of the two.

Because $+3$ therefore sits *inside* the noise band of any single run, the protocol never let it stand alone. A Go verdict additionally requires that at least two of the three seeds in a cell be positive, that GRPO beat RS-SFT rather than only plain SFT, and that the verifier-hacking audit be clean; and an effect anywhere between two and five points, or one whose seeds disagree in sign, is declared INCONCLUSIVE and sent back for two further seeds rather than adjudicated. The evidence this study rests on is accordingly agreement across seeds and cells rather than per-group significance—which was the design, not a retreat, as Section 8 quantifies.

Solver cost. The solver is task-specific engineering, requiring roughly 440 new lines of adaptation code on top of a reusable symbolic execution core. We report this because the fair reading of our result is not “a small script beats RL” but “a few hundred lines of task-specific symbolic code, which this benchmark family admits by construction, beats RL under scarcity.”

5 Results

5.1 Every frozen-test effect size is negative

Table 1 and Figure 2 give the small-backbone result: 11 groups, Δ from -1.42 to -16.42 points, mean -5.48 . No group approaches the pre-registered $+3$ -point threshold. Adding the six 7B groups of Section 6 gives 17 negative values out of 17, mean -4.44 .

The two configurations differ in dataset family, model family, and absolute difficulty—the solver scores 64.54% on Config 1 and 32.08–32.83% on Config 2—yet agree in sign. Config 2 at $m=128$ is significant under McNemar after Holm correction ($p < 0.001$).

One group deserves note for a reason unrelated to Δ : at Config 2 $m=128$, GRPO (16.11–17.02%) falls below continued SFT (18.83–20.18%) and RS-SFT (21.84–25.00%). Here the verdict is negative even without the solver—the learning-only control group already suffices.

5.2 Why: the scarce regime is the non-learner’s regime

Figure 3 explains the sign at small scale. Warm-start accuracy in Config 1 climbs from 0.0% at $m=0$ through 19.0% ($m=2$), 33.0% ($m=64$) and 65.8% ($m=256$) to 80.2% at the full pool; Config 2 climbs from 0.5%

Table 1: Frozen-test accuracy (%) per arm and Δ (points) at 0.5B and 1B. The maximum over $\mathcal{B}$ is attained by the solver in all 11 groups. Config 2 $m=128$ seed 1 is excluded: its GRPO run ended with zero positive reward (Section 7).

m	seed	solver	W_m	c-SFT	RS-SFT	GRPO	Δ
<i>Config 1: CodeTAT-QA $\times$ Qwen2.5-0.5B ($n=282$)</i>							
2	0	64.54	15.60	24.47	27.66	62.41	−2.13
2	1	64.54	14.54	18.09	28.01	63.12	−1.42
2	2	64.54	26.60	20.57	30.50	60.64	−3.90
64	0	64.54	29.43	43.62	51.06	58.51	−6.03
64	1	64.54	27.66	45.39	56.03	60.99	−3.55
64	2	64.54	32.62	45.74	51.42	61.70	−2.84
<i>Config 2: CodeFinQA $\times$ Llama-3.2-1B ($n=664$)</i>							
128	0	32.38	12.05	20.18	21.84	17.02	−15.36
128	2	32.53	12.65	18.83	25.00	16.11	−16.42
256	0	32.38	25.60	26.36	28.92	29.97	−2.41
256	1	32.83	25.75	27.11	25.90	28.16	−4.67
256	2	32.08	28.16	27.41	30.27	30.57	−1.51

to 47.0%. The solver consumes no parameters and learns nothing: it scores 72.0% at *every one* of Config 1’s ten budgets, and stays within 34.5–38.0% across Config 2’s seven.

The crossing points matter. In Config 1 the learning curve does not overtake the solver until $m=1024$ (74.5% vs. 72.0%); at $m=256$ it is still behind at 65.8%. Config 2 crosses near $m=1024$ as well (43.0% vs. 36.3%). Everything left of that crossing is the demonstration-scarce regime—the regime in which RLVR is most often argued to be valuable, and in which a method that ignores demonstrations entirely is hardest to beat.

5.3 The gain is real; the conclusion is not

Against a warm-start-only control the same runs look excellent: on dev, GRPO improves over W_m from 19.0% to 72.0% at $m=2$ and from 33.0% to 67.0% at $m=64$. Both would be reported as large positive results. Adding one non-learning arm to the same comparison turns every frozen-test verdict negative. RL is doing real work here—it beats every learning baseline by up to 35 points—but “does RL beat SFT?” and “is RL the best use of this demonstration budget?” have different answers.

5.4 Dev selection bias, observed three times

On dev, two groups produced non-negative values: Config 1 $m=2$ gave $\Delta = +0.00, -0.50, +0.50$. The same three groups on the frozen test gave $-2.13, -1.42, -3.90$. Two further instances appear later: the $\beta=0$ hyperparameter setting (Section 7) yields $+1.50$ on dev and -2.13 on test, and the 7B $m=2$ seed-2 group yields $+1.50$ on dev and -1.77 on test. In all three, dev favors GRPO and the frozen test does not.

Since every checkpoint in every arm is selected on dev, the arm whose dev trajectory fluctuates most gains the most from that selection—and a policy-gradient run plausibly fluctuates more than a fine-tuning run. We state this as a candidate mechanism rather than a measured one: three instances identified after the fact do not establish it, and our stored artifacts support the per-arm dev-checkpoint variance it predicts only at 7B (Section 9). What the three instances do show without interpretation is the practical point: without a frozen test set we would have reported non-negative results in three groups and a positive one from the sweep. All $17 \times 5 = 85$ arm evaluations of the reported groups are present; nothing in $\mathcal{B}$ is imputed or dropped.

Table 2: Frozen-test accuracy (%) at 7B, all three seeds. The solver is now 15 points behind and irrelevant; continued SFT attains the maximum in all six groups and GRPO trails it in all six.

m	seed	solver	W_m	c-SFT	RS-SFT	GRPO	Δ
2	0	64.54	71.28	79.79	75.18	76.60	−3.19
2	1	64.54	69.86	77.66	76.95	74.82	−2.84
2	2	64.54	70.57	76.95	74.82	75.18	−1.77
64	0	64.54	79.08	81.91	80.50	79.79	−2.13
64	1	64.54	78.72	80.50	80.14	79.43	−1.06
64	2	64.54	79.43	85.11	81.21	80.85	−4.26

6 Scale: the winning baseline changes identity

The standing objection to everything above is that 0.5B and 1B models are small enough that RLVR is not expected to help, making the negative result unsurprising. We test that objection rather than concede it: the same CodeTAT-QA task, the same protocol, the same two budgets, three seeds, on Qwen2.5-7B-Instruct.

Table 2 shows the transition. The solver’s accuracy is unchanged by construction—64.54%, the same number as at 0.5B, because it contains no model—but the learning arms have moved past it. Warm-start SFT on two demonstrations now reaches 69.86–71.28%, and every arm at $m=64$ is around 80%. The crossing point predicted by the dose-response curves of Section 5 is real and lies between 1B and 7B on this task. The “gains emerge with capability” account is correct about the solver.

It is not correct about Δ . In all six groups the maximum is attained by **continued SFT**—76.95–85.11%—and GRPO trails it by 1.06 to 4.26 points (mean −2.54). Scaling the backbone 14× replaced the arm that beats GRPO; it did not produce a setting in which GRPO leads.

What this does to the claim. Within the registered budget it generalizes it. Had the solver kept winning at 7B, a reader could reasonably suspect we had over-engineered one baseline for one regime. Instead the solver exits exactly where the scaling account says it should, and the arm that replaces it is continued SFT—an arm present in essentially every RLVR paper, requiring no engineering, and impossible to characterize as an exotic construction. The methodological requirement that survives both regimes is not “include a solver” but “make the control group broad enough that you would notice, and aggregate it with a max.”

But this configuration does not answer the question it was added for, and we should say so before a reader finds out. Config 3 exists to test the objection that 0.5B and 1B are simply too small. It cannot settle that objection, because at 7B the registered wall-clock budget buys an order of magnitude less data than it does at 0.5B: 0.13 epochs over the prompt pool against 0.53–0.79. We re-ran this section’s smallest-margin cell— $m=64$, seed 1, the one most favourable to GRPO—at ten times the coverage, and it reverses decisively, to $\Delta = +6.74$ and $+6.38$ across two training seeds, clearing the pre-registered Go threshold (Appendix G). At the registered budget the result holds, including under a run that removes the zero-advantage degeneracy we diagnose in Section 7. **What Config 3 demonstrates is therefore a budget boundary, not a capability one.** We keep it in the paper for that reason—knowing where a null result stops is worth more than a scale ladder that does not hold—but the quantitative weight of this study rests on the 11 groups at 0.5B and 1B, not on these six.

Statistical caveats. Effect sizes at 7B are smaller than at 0.5B/1B: −1.06 to −4.26 points on $n=282$ is 3 to 12 items, and the three learning arms lie within a few points of one another, so which of them attains $\max \mathcal{B}$ is itself sensitive to sampling noise. Within the registered budget the 7B groups accordingly carry

Table 3: Solver ablation (dev, mean of 3 seeds). Retrieval never fires: its marginal contribution $C - A$ is +0.00 points at all ten Config 1 budgets.

m	A: rules	B: retrieval	C: deployed	$C - A$
<i>Config 1 (all 10 budgets identical)</i>				
0 ... 2446	72.00	0.00	72.00	+0.00
<i>Config 2</i>				
128	36.00	0.00	35.33	−0.67
256	35.67	0.00	35.33	−0.33

a qualitative claim—the arm that beats GRPO changes identity, and GRPO is not that arm in any of the six—rather than an effect size, and past that budget they do not carry even that. The quantitative result rests on the 11 groups at 0.5B and 1B, where $\max \mathcal{B}$ is attained by a deterministic solver that leads the nearest learning arm by *at least* 8.5 points in Config 1 (8.51 to 36.88 points across the six groups; the quoted figure is the smallest, not the typical): there the identity of the maximum is not in question, and no selection over arms enters the estimate.

The largest 7B margin (−4.26 points, Config 3 $m=64$ seed 2) has an uncorrected McNemar p of 0.017. Under the Holm correction this paper declares across its family of 17 groups, that value becomes 0.254 and is **not** significant; the conference version of this paper quoted the uncorrected figure alongside a stated Holm procedure, which was inconsistent. No 7B group is significant after correction. The 7B claim is the qualitative one stated above, supported by six of six groups agreeing in sign (Section 8), and by nothing stronger.

7 Ruling Out the Obvious Objections

A negative result of this shape invites four specific rebuttals: the solver is memorizing, its strength is generic to any non-learning method, GRPO’s score is a reward-hacking artifact, and GRPO is simply mistuned. Each is testable; none survives. We then report the one run we excluded, ask whether our GRPO is simply under-trained, and state the solver’s cost.

7.1 “The solver is just retrieving gold answers”

If the solver’s accuracy came from copying gold programs out of its retrieval library, it would be a memorization artifact rather than a baseline. We decompose it into three variants: **A** uses only the hand-written rule path; **B** uses only retrieval; **C** is the deployed configuration (rules first, retrieval as fallback).

Table 3 shows the retrieval-only variant scores 0.00% everywhere: it never produces an accepted program. Across all ten Config 1 budgets and three seeds—thirty measurements— A , C and their difference are constant at 72.00%, 72.00% and +0.00 points. In Config 2 the deployed variant differs from the rule-only path by −0.33 to −0.67 points—one to two items on the 200-item dev subset, i.e. within noise—so retrieval contributes nothing measurable on either configuration. The solver’s ability is therefore entirely symbolic. This also explains the flat dose curve: with retrieval never firing, the demonstration budget has no channel through which to affect the solver at all.

7.2 “Any non-learning baseline would look that strong”

A second objection is that the solver’s advantage demonstrates nothing about symbolic competence: on a task whose answers are executable programs, perhaps any method that skips training lands near the same

Table 4: Frozen-test accuracy (%) of the two non-learning options. Prompting the base model directly is by far the weakest arm in both configurations, so the solver’s advantage does not come free with skipping training.

	Config 1	Config 2
solver (symbolic rules)	64.54	32.08–32.83
PoT / PAL prompting (no training)	9.93–16.67	0.75–1.36
GRPO (RL)	58.51–63.12	16.11–30.57

Table 5: Verifier-hacking audit over the 14 GRPO runs whose reward dumps were retained. H1: constant program; H4: no-op arithmetic; H3: share of positive samples under the single most frequent code skeleton.

	samples	positive	H1	H4	H3
Config 1 (6 runs)	87,872	47,616	0	4	$\leq 0.9\%$
Config 2 (6 runs)	37,933	6,030	0	23	$\leq 1.5\%$
Config 3, 7B (2 runs)	4,672	3,871	0	0	$\leq 2.1\%$

place. The natural test is the other obvious non-learning option: prompt the base model directly to emit a program, with the same m demonstrations used as few-shot examples. This is Program-of-Thought / PAL prompting [21, 22]—no gradient steps at all.

Table 4 gives the result over the eleven groups of Table 1. Few-shot prompting reaches 13.30% on average in Config 1 and 1.00% in Config 2—the weakest arm in the entire study, far below both the solver and GRPO. Two things follow. First, the solver’s strength is specific rather than generic: the most obvious alternative non-learning baseline is dramatically worse, so what the solver contributes is genuine symbolic competence rather than a free advantage of skipping training. Second, RLVR is doing substantial work—GRPO improves over prompting the same base model by 44 to 50 points across the Config 1 groups—which is precisely why the comparison matters: a method that clearly helps can still be the wrong way to spend the budget.

7.3 “GRPO’s score is inflated by verifier hacking”

If GRPO were exploiting the reward checker—emitting constant programs, inlining the gold value, or collapsing onto one template—its accuracy would be an artifact and the comparison meaningless. We audit reward-dump records across the 14 GRPO runs whose reward dumps were retained, spanning all three backbones.

Table 5 reports the result. Constant-program shortcuts occur **zero** times in 57,517 positive-reward samples—the verifier’s anti-constant guard holds. No-op arithmetic appears 27 times (0.05%). Template collapse is absent: the most frequent code skeleton accounts for at most 2.1% of positive samples, with skeleton Gini coefficients between 0.06 and 0.45, indicating a broadly distributed solution space. Gold-literal matches are rare (0.19%) and expected on these tasks, where gold answers are computed from context numbers.

GRPO’s 16.1–80.9% frozen-test accuracies across the three backbones therefore reflect genuine program synthesis. That is precisely what makes the comparison meaningful: a policy that is really solving the task still loses to a solver that never learned.

7.4 “You did not tune GRPO properly”

A fourth objection is that the negative result holds only at our particular hyperparameters. We sweep the knob most likely to matter—the KL coefficient β —at Config 1’s higher operating budget, in its smallest-margin

Table 6: GRPO hyperparameter sweep, Config 1 $m=64$, seed 2 (solver: 72.00 dev / 64.54 test). $\beta=0$ is the only setting whose Δ turns positive—and only on dev.

β	temp	dev	test	Δ_{dev}	Δ_{test}
0.04	1.0	68.00	61.70	−4.00	−2.84
0.0	1.0	73.50	62.41	+1.50	−2.13
0.2	1.0	58.00	51.42	−14.00	−13.12

Table 7: Positive-reward rate across all 18 GRPO runs. One run, on the 1B backbone at the scarcest budget, ends at zero and is excluded.

Backbone	budget	runs	positive-reward rate
Qwen2.5-0.5B	2, 64	6	0.522–0.560
Llama-3.2-1B	128	3	0.081, 0.092, 0.000
Llama-3.2-1B	256	3	0.194–0.239
Qwen2.5-7B	2, 64	6	0.772–0.846

seed ($m=64$, seed 2), and evaluate every resulting policy on both dev and the frozen test set.

Table 6 contains the sharpest single observation in this paper. Removing the KL penalty entirely ($\beta=0$) does lift GRPO above the solver *on dev*, by 1.50 points—the only positive Δ any setting produced. On the frozen test set the same policy scores 62.41% against the solver’s 64.54%: $\Delta = -2.13$ points. The reversal is a selection artifact, not an effect.

This is one of three independent instances of the same phenomenon (Section 5): Config 1 at $m=2$ gave +0.50 on dev against −3.90 on the frozen test, and the 7B $m=2$ seed-2 group gave +1.50 against −1.77. In all three, dev favored RL and the frozen test did not. Had we followed the common practice of reporting dev numbers after a hyperparameter search, we would have published the opposite conclusion—from the same runs.

7.5 One excluded run

One group (Config 2, $m = 128$, seed 1) experienced training collapse, ending with a zero positive-reward rate. We exclude this run as an infrastructure/training failure rather than an algorithmic property; notably, including it ($\Delta = -19.73$) would further strengthen our negative finding.

Table 7 places it in context: the collapse is confined to the weakest backbone at the scarcest budget, and every other run in the study sustains a positive reward rate throughout training.

7.6 Is our GRPO simply under-trained?

At 7B, GRPO ranks third of five arms on seed-averaged accuracy, behind continued SFT and RS-SFT. Since beating SFT is RLVR’s core claim, a reader should ask whether our GRPO is broken rather than our conclusion interesting. Four pieces of evidence say it is not.

First, *it works where it should*: at 0.5B, GRPO beats every learning arm by wide margins—62.41% against warm-start’s 15.60% and continued SFT’s 24.47% at $m=2$. A broken implementation would not produce a 35-point gain in one regime and a 3-point deficit in another.

Second, *the pre-registered budget was consumed*: every GRPO run reached its 5700-second cap with 2240–17920 verifier calls and 131k–947k generated tokens. Two qualifications are needed here, both of which the conference version of this paper omitted.

Table 8: Compute on four axes, ranges over the three seeds. Wall clock is the pre-registered matching axis; the other three are reported so that the matching is not taken on trust. RS-SFT, not GRPO, is the most expensive arm on generated tokens and verifier calls.

cfg	m	wall clock (s)			verifier calls		generated tokens	
		c-SFT	GRPO	RS-SFT	GRPO	RS-SFT	GRPO	RS-SFT
C1	2	2579–3196	5701–5703	5715–5732	15488–17920	21824–24576	722k–798k	5441k–6136k
C1	64	5703–5707	5711–5728	4561–5815	12160–13184	13184–15872	908k–947k	2906k–3301k
C2	128	5704–5715	5709–5751	1549–5905	6272–7040	7040–9728	296k–323k	1103k–1298k
C2	256	5703–5714	5720–5755	5822–6038	7104–7872	6144–6656	299k–335k	602k–762k
C3	2	5702–5706	5712–5761	—	2304–2880	—	131k–150k	—
C3	64	5721–5748	5750–5864	—	2240–2496	—	167k–181k	—

The first concerns what “matched compute” means. GRPO receives the same wall clock as continued SFT, which is the arm attaining $\max \mathcal{B}$ in all six 7B groups, and vastly more than the solver, which attains it in the other eleven and runs on CPU in seconds. But GRPO does *not* exceed every arm on every axis: RS-SFT generates three to seven times more tokens than GRPO in Config 1 (Table 8). RS-SFT never attains $\max \mathcal{B}$ in any of the 17 groups, so this does not enter Δ —but the blanket claim that GRPO received at least as much compute as the arms it is compared with is true only on the wall-clock axis and only against the arms that actually bind.

The second qualification is more serious, and it is the strongest limitation in this study. Each prompt is sampled 8 times, so the verifier-call counts above imply that GRPO visited 280–2240 distinct prompts. Against pools of 2846 and 4035 items, **no GRPO run in this study completed even one epoch**:

config	m	verifier calls	distinct prompts	epochs over pool
Config 1 (0.5B)	2	15488–17920	1936–2240	0.68–0.79
Config 1 (0.5B)	64	12160–13184	1520–1648	0.53–0.58
Config 2 (1B)	128	6272–7040	784–880	0.19–0.22
Config 2 (1B)	256	7104–7872	888–984	0.22–0.24
Config 3 (7B)	2	2304–2880	288–360	0.10–0.13
Config 3 (7B)	64	2240–2496	280–312	0.10–0.11

Coverage is worst exactly where the paper’s generalization claim lives: at 7B, GRPO saw roughly a tenth of the available prompts. Equal wall clock at a fixed per-token cost buys an order of magnitude less data at 7B than at 0.5B, and our pre-registration fixed the seconds rather than the data. We do not think this explains the result—the argument continues below—but a reader who believes that RLVR needs multiple epochs to pay off has not been answered by this study, and we do not claim otherwise.

Third, *the policy is genuinely solving*: the audit of Table 5 covers the 7B runs directly and finds zero constant-program shortcuts, zero no-op arithmetic, and a top-skeleton share of 2.1% over 3871 positive samples.

The conference version of this paper continued the third argument by noting that the 7B positive-reward rate, 0.772–0.846, is the highest of the three backbones, and read that as evidence that the 7B policy was training successfully. **That inference is backwards, and we retract it.** GRPO’s advantage is computed within a group of samples drawn from the same prompt; if all eight samples receive the same reward the group’s advantage is zero and it contributes no gradient. A high positive-reward rate therefore implies *more* degenerate groups, not more learning.

Table 9 makes the size of the effect visible. At 7B, between 70.7% and 86.1% of prompt groups are all-

Table 9: Zero-advantage groups across all 18 GRPO runs, read from each run’s recorded metadata. A prompt group whose 8 samples all receive the same reward has zero group-relative advantage and contributes no gradient. The 7B runs have the highest positive-reward rate in the study and the fewest informative prompts.

backbone	m /seed	steps	groups	all-0	all-1	zero-var	informative	pos. rate
Qwen2.5-0.5B	2/0	242	1936	565	652	62.9%	719	0.527
	2/1	280	2240	625	871	66.8%	744	0.560
	2/2	258	2064	572	748	64.0%	744	0.551
	64/0	190	1520	381	387	50.5%	752	0.522
	64/1	197	1576	348	427	49.2%	801	0.544
	64/2	206	1648	386	441	50.2%	821	0.540
Llama-3.2-1B	128/0	98	784	544	1	69.5%	239	0.081
	128/1 [†]	110	880	880	0	100.0%	0	0.000
	128/2	102	816	549	2	67.5%	265	0.092
	256/0	111	888	463	16	53.9%	409	0.196
	256/1	111	888	461	12	53.3%	415	0.194
	256/2	123	984	460	36	50.4%	488	0.239
Qwen2.5-7B	2/0	38	304	29	207	77.6%	68	0.812
	2/1	36	288	42	185	78.8%	61	0.772
	2/2	45	360	43	267	86.1%	50	0.823
	64/0	35	280	18	180	70.7%	82	0.846
	64/1	39	312	24	206	73.7%	82	0.837
	64/2	35	280	21	194	76.8%	65	0.826

[†]The excluded run: every group all-zero.

correct or all-wrong, leaving **50 to 82 informative prompts for an entire run**. At 0.5B the same quantity is 719 to 821. The 7B runs are thus training on roughly an order of magnitude less gradient signal than the 0.5B runs that this study reads as its positive control. The excluded Config 2 run appears here in its proper form: 880 of 880 groups all-zero, zero informative prompts.

The honest position is that the third argument survives only in its first half. The 7B policy is solving the task rather than gaming the verifier, and that much the audit establishes. It is *not* established that the 7B GRPO arm received enough gradient signal to represent what GRPO can do at that scale.

We subsequently tested that directly rather than leaving it as a caveat, and the answer is that the 7B arm was indeed under-trained: at ten times the coverage the smallest-margin 7B cell reverses from -1.06 to $+6.74$ (Appendix G). Two things about that result belong here. It is a *budget* effect and not a degeneracy effect—a matched-budget run that raises informative prompts from 82 to 120 moves Δ to $+1.06$ with an interval of $[-2.13, +4.26]$, which is the same place as the original. And the degeneracy does not improve with training: across $10.6\times$ the steps the zero-advantage fraction *rose* from 73.7% to 86.0%, so the win came from brute force rather than from better use of the signal. Section 6 states what this costs the 7B configuration.

Fourth, *tuning does not rescue it*: the sweep in Table 6 covers the KL coefficient. The single setting that lifts GRPO above the best baseline does so only on dev, and reverses on the frozen test.

The more economical reading is that a stronger backbone helps plain imitation at least as much as it helps RL. At 7B, continued SFT on m demonstrations is already close to what the model can do on this task; the room RL has to add value has narrowed, while its cost has not.

7.7 “You left out the obvious inference-time baseline”

A fifth objection, which we did not anticipate and which reviewers of the conference version raised independently, is that the control group omits the cheapest non-learning option of all: sample n completions from the warm start and keep the one the verifier accepts. If verifier-filtered best-of- n were strong, it would be a far more natural control arm than a hand-written solver.

It is not available on these tasks, for a structural reason. Our verifier decides correctness by executing the emitted program and comparing its output to the *gold* answer within a 1% relative tolerance. At training time that is legitimate—the gold answer is part of the training item. At inference time it is not: selecting among n candidates with this verifier requires knowing the answer. Verifier-filtered best-of- n on these tasks is therefore identical to $\text{pass}@n$, an oracle upper bound rather than a deployable arm. The deployable relative is majority voting over the n samples, which needs no gold.

We measured both from the warm start at $n=16$ on dev:

config	m	pass@1	pass@16 (oracle)	maj@16 (deployable)	solver
Config 1	2	9.83	48.83	35.33	72.00
Config 1	64	13.17	61.17	45.17	72.00
Config 2	128	5.50	34.50	17.17	35.33
Config 2	256	16.00	51.67	30.67	35.33

At all four operating points the deployable arm (maj@16) is below the solver, by 15 to 37 points. At three of the four, so is the *oracle* upper bound: even an omniscient selector over 16 samples from the warm start would not reach the solver at Config 1 $m=2$, Config 1 $m=64$, or Config 2 $m=128$. Only Config 2 $m=256$ has oracle headroom above the solver, and there the deployable version still falls 4.7 points short.

Two caveats. These are dev numbers on the 200-item selection subset, not frozen test; and this arm was not part of the pre-registered control group, so it enters as supporting evidence rather than into Δ . Appendix C gives the full budget sweep, in which the sampling arms do overtake the solver past $m=256$ —the same crossing the dose curves show for every learning arm.

7.8 How the solver was built, and what it cost

Because the solver attains $\max \mathcal{B}$ in 11 of 17 groups, its construction is the single largest discretionary choice in this study, and a reader is entitled to the procedure rather than an assurance. The pre-registered protocol constrained the solver in one direction only—it forbade weakening it if GRPO won—which does nothing to prevent the opposite failure. What actually constrains it is where it was allowed to look.

The solver was developed against a dev split carved from training data (for Config 2, the self-consistent subset of `train at idx % 4 == 0`, $n=1013$), with leave-one-out retrieval, under a fast in-process replica of the verifier’s semantics. The frozen test set was evaluated *once*, at the end, with the real verifier. Development proceeded over numbered rounds whose dev accuracy we logged as it moved— $0.243 \rightarrow 0.294 \rightarrow 0.317 \rightarrow \dots$ for Config 2—together with the failure mode each round targeted. Those logs are published verbatim (Appendix F); they are the audit trail for this arm, and they are not flattering in places.

The cost of that dev fitting is visible and we should name it. The solver scores 72.00% on Config 1 dev and 64.54% on Config 1 frozen test: a drop of 7.46 points, larger than 15 of the 17 effect sizes this paper reports. That gap is the solver’s own generalization penalty from having been iterated against dev, and it is already fully paid inside every number in Table 1, because Δ is computed on test. A reader who suspects the solver of being dev-overfitted is right, and the frozen-test protocol has already charged it for that.

The engineering cost was roughly 440 lines of task-specific adaptation code on top of a reusable symbolic execution core, developed over the logged rounds, after which inference runs in seconds on CPU. We

Table 10: Solver transferability. Each solver is run unmodified on the other task’s frozen test set, under the same clean-subset protocol. The narrative solver transfers and still beats all three learning arms of Config 1 at $m=2$; the table solver does not transfer at all, because it cannot parse a context that is not a JSON table and abstains on every item.

solver	home task	other task	retained
built on CodeTAT-QA	64.54	0.00	0%
built on CodeFinQA	33.43	36.17	108%

do not claim this is cheap in human terms, and “simpler” in this paper’s title refers to the method’s inference-time and data requirements, not to the labour of writing it. Whether a practitioner should build such an arm is a different question from whether a paper claiming an RLVR gain should have to clear it.

Does the solver transfer? The sharpest form of the objection is that 440 dataset-specific lines are not “simpler” in any useful sense, because they buy nothing outside the dataset they were written for. This is testable within our own study: we have two tasks and two independently developed solvers, so each can be run, unmodified, on the other’s frozen test set under the same clean-subset protocol.

Table 10 shows that the answer depends on direction, and that the objection is half right. The CodeTAT-QA solver does not transfer at all: it parses a JSON table out of the context and emits cell-indexed programs, and on CodeFinQA’s narrative contexts it finds no table and abstains on all 664 items. That is an architectural incompatibility rather than a graceful degradation.

The CodeFinQA solver transfers. Run unmodified on CodeTAT-QA it scores 36.17%, slightly *above* its 33.43% on its home task. The comparison that matters is with Table 1: at Config 1’s $m=2$ operating point the three learning arms score 14.54–26.60% (W_m), 18.09–24.47% (c-SFT) and 27.66–30.50% (RS-SFT). A solver written for a different dataset, carried across without a line changed, beats all three. It does not approach the native solver’s 64.54%, and it would not have attained $\max \mathcal{B}$ in any reported group—but it establishes that the control-group requirement does not depend on bespoke engineering for the exact benchmark under test. Half of a symbolic baseline is portable; the table parser is not.

8 The Decision Quantity, Audited

Everything in this section is computed from the stored per-item outcome vectors of the same runs reported above, by a single script (`analysis/gen_extended_stats.py`); no value is transcribed by hand. The section exists because a decision rule built on a max invites two objections that the conference version had no room to answer, and because a study reporting a null owes the reader its per-group statistics rather than the two cells that happen to be significant.

8.1 The pre-registered question: does GRPO clear +3?

The decision rule this study registered is a Go threshold, not a test of inferiority: GRPO passes if $\Delta \geq +3$ points. That is the hypothesis the analysis owes the reader an answer to, and it is a weaker and more answerable one than “is GRPO significantly worse.” Conflating the two understates the result, so we report both, starting with the registered one.

A paired item bootstrap ($B = 50,000$, max recomputed on each resample) gives, per group, the probability that Δ reaches the Go threshold:

Table 11: Per-group significance and power. b/c are the McNemar discordant counts (GRPO right / arg max $\mathcal{B}$ right). p_{Holm} corrects across all 17 groups. The CI is a paired item bootstrap ($B=10,000$) of Δ with the max recomputed on each resample. MDE is the smallest $|\Delta|$ this group’s discordant count could have declared significant, uncorrected and under Holm.

m	seed	arg max $\mathcal{B}$	Δ	b/c	p	p_{Holm}	95% CI	MDE	MDE _H
<i>Config 1 (CodeTAT-QA $\times$ 0.5B, $n=282$)</i>									
2	0	solver	-2.13	38/44	0.581	1.000	$[-8.51, +3.90]$	7.09	9.93
2	1	solver	-1.42	37/41	0.734	1.000	$[-7.45, +4.61]$	7.09	9.93
2	2	solver	-3.90	32/43	0.248	1.000	$[-9.93, +2.13]$	6.74	9.57
64	0	solver	-6.03	30/47	0.068	0.878	$[-12.06, +0.00]$	6.74	9.57
64	1	solver	-3.55	30/40	0.282	1.000	$[-9.22, +2.13]$	6.38	9.22
64	2	solver	-2.84	32/40	0.410	1.000	$[-8.87, +3.19]$	6.38	9.93
<i>Config 2 (CodeFinQA $\times$ 1B, $n=664$)</i>									
128	0	solver	-15.36	60/162	5.1×10^{-12}	8.2×10^{-11}	$[-19.58, -11.14]$	4.82	6.93
128	2	solver	-16.42	59/168	2.6×10^{-13}	4.4×10^{-12}	$[-20.63, -12.05]$	4.67	7.08
256	0	solver	-2.41	110/126	0.329	1.000	$[-6.93, +1.51]$	4.82	7.23
256	1	solver	-4.67	88/119	0.037	0.515	$[-9.04, -0.45]$	4.67	6.78
256	2	solver	-1.51	105/115	0.544	1.000	$[-5.87, +1.66]$	4.82	6.93
<i>Config 3 (CodeTAT-QA $\times$ 7B, $n=282$)</i>									
2	0	c-SFT	-3.19	6/15	0.078	0.940	$[-6.38, +0.00]$	3.90	5.32
2	1	c-SFT	-2.84	9/17	0.169	1.000	$[-6.38, -0.35]$	4.26	5.67
2	2	c-SFT	-1.77	15/20	0.500	1.000	$[-6.03, +1.42]$	4.61	6.74
64	0	c-SFT	-2.13	9/15	0.307	1.000	$[-5.32, +0.00]$	4.26	5.67
64	1	c-SFT	-1.06	10/13	0.678	1.000	$[-4.26, +0.71]$	3.90	5.32
64	2	c-SFT	-4.26	5/17	0.017	0.254	$[-7.45, -1.06]$	4.26	5.67

question	criterion	groups
Does GRPO clear the registered +3 threshold?	95% CI upper bound $< +3$	14/17
Is GRPO strictly worse than max $\mathcal{B}$?	95% CI upper bound < 0	5/17

$P(\Delta \geq +3)$ has a median of 0.0013 across the 17 groups and does not exceed 0.0804 in any of them. The three groups whose intervals fail to exclude +3 are all in Config 1— $m=2$ seeds 0 and 1, and $m=64$ seed 2—where $n=282$ and GRPO and the solver disagree on 70 to 82 items, giving a wide paired interval. Config 3 has the same n but excludes +3 in all six groups, because GRPO and continued SFT agree on all but 21 to 35 items and the paired interval is correspondingly narrow. Counts at the +3 boundary move by one group between bootstrap runs; the continuous $P(\Delta \geq +3)$ does not, and should be preferred.

So the registered claim—GRPO does not reach the threshold that would have made this a positive result—is supported in 14 of 17 groups, and the remaining three are inconclusive rather than contrary. The stricter claim that GRPO is *worse* is supported in 5. We make the first claim, not the second.

8.2 Per-group significance, and what it does not show

Table 11 reports, for all 17 groups, the McNemar discordant counts, the exact two-sided p , the same p under Holm correction across the full family of 17, and a paired item bootstrap interval for Δ . These test the stricter hypothesis, and the reader should keep Section 8.1 in view while reading them.

The honest summary is unflattering. Under Holm correction **only two groups are significant**, both of them Config 2 at $m=128$ ($p_{\text{Holm}} = 8.2 \times 10^{-11}$ and 4.4×10^{-12}). The bootstrap interval excludes zero in five

groups. Removing the two Config 2 $m=128$ groups—the cell where one seed also collapsed—the remaining 15 groups average -2.91 points and not one of them is individually significant. A reader who requires per-group significance should read this paper as establishing a negative result on one cell and a consistent sign everywhere else.

That one cell deserves an additional warning against itself. Config 2 at $m=128$ is not merely the cell with the largest effect; it is the cell in which the RL arm was least healthy. Its two surviving runs have positive-reward rates of 0.081 and 0.092, the lowest of any run in the study, its third seed collapsed to zero entirely, and its prompt-pool coverage is 0.19–0.22 epochs. The only two groups that survive multiple-comparison correction are therefore the two in which under-training is the most plausible competing explanation. We do not think this invalidates them—continued SFT and RS-SFT beat GRPO there without the solver’s help (Table 15), so the verdict does not rest on the solver—but a reader should not treat those two p -values as the study’s strongest evidence. The strongest evidence is the Go-threshold result of Section 8.1 and the sign pattern below.

The last two columns explain why. Given each group’s discordant count, the smallest $|\Delta|$ that its McNemar test could have declared significant is a median of 4.82 points uncorrected and 6.93 points under Holm. The median $|\Delta|$ we actually measure is 2.84 points, and the pre-registered Go threshold is $+3$. *No single group in this study can resolve a three-point effect in either direction.*

We want to be exact about what that does and does not concede. It is not a gap the analysis discovered after the fact: the protocol fixed $+3$ precisely *because* it put single-seed standard error at 3–5 points (Section 4), so the threshold was knowingly set at the floor of the design’s resolution, and the Go rule was built to compensate—two of three seeds positive, an INCONCLUSIVE band over 2–5 points, further seeds rather than a verdict when signs disagree. The numbers above confirm that estimate rather than contradicting it: a median MDE of 4.82 is what an SE of 3–5 predicts.

What it does concede is symmetric and real. Had GRPO enjoyed a genuine $+3$ -point advantage, this design could not have certified it group by group either; it would have returned INCONCLUSIVE and demanded more seeds. Per-group significance was therefore never the evidence, by construction. The sign pattern is.

8.3 What the sign pattern is worth

All 17 groups are negative. Treating them as independent observations gives a sign-test $p = 1.5 \times 10^{-5}$, but they are not independent: three seeds within a cell share a test set, a solver and a task, and the six Config 1 and Config 3 cells share the test set entirely. We therefore report the same test at three levels of clustering:

unit of independence	units	sign-test p
group (3 seeds $\times$ 2 budgets $\times$ 3 configs, most optimistic)	17	1.5×10^{-5}
(config, budget) cell	6	0.031
configuration	3	0.250

The defensible reading is the middle row: six independent operating points, all negative, $p = 0.031$. The headline “17 negative effect sizes” is a description of the data, not seventeen pieces of evidence, and we do not claim it as such.

8.4 Is the `max` itself producing the sign?

The sharpest objection to this design is that it commits the error the paper was written to criticize. Δ subtracts a `max` over four arms, that `max` is taken on the frozen test set, and a maximum over noisy estimates is upward-biased—so Δ is biased downward under the null, by an amount that grows with the number of arms

Table 12: Winner’s-curse audit. *bias* is the bootstrap mean of Δ with the max recomputed on each resample minus the same quantity with the arm frozen at the full-sample arg max. Δ_{cf} is cross-fitted: the control arm is selected on a random half of the test set and scored on the other half ($K=2,000$ splits).

m	seed	Δ (naive)	bias	Δ_{cf}	95% CI
<i>Config 1 (CodeTAT-QA $\times$ 0.5B, $n=282$)</i>					
2	0	−2.13	+0.000	−2.29	[−8.51, +4.26]
2	1	−1.42	+0.000	−1.41	[−7.80, +4.96]
2	2	−3.90	+0.000	−3.89	[−9.93, +2.13]
64	0	−6.03	+0.000	−6.02	[−12.06, +0.71]
64	1	−3.55	−0.002	−3.55	[−9.22, +2.13]
64	2	−2.84	+0.000	−2.86	[−9.22, +2.84]
<i>Config 2 (CodeFinQA $\times$ 1B, $n=664$)</i>					
128	0	−15.36	+0.000	−15.29	[−19.28, −11.14]
128	2	−16.42	−0.000	−16.40	[−20.78, −12.05]
256	0	−2.41	−0.063	−1.99	[−6.33, +2.71]
256	1	−4.67	−0.004	−4.64	[−8.73, −0.30]
256	2	−1.51	−0.274	−0.50	[−3.92, +3.31]
<i>Config 3 (CodeTAT-QA $\times$ 7B, $n=282$)</i>					
2	0	−3.19	−0.002	−3.12	[−6.38, +0.00]
2	1	−2.84	−0.364	−2.10	[−4.96, +0.71]
2	2	−1.77	−0.172	−1.08	[−4.96, +2.84]
64	0	−2.13	−0.200	−1.40	[−4.26, +1.42]
64	1	−1.06	−0.482	−0.11	[−2.84, +2.84]
64	2	−4.26	−0.003	−4.18	[−7.09, −0.71]

and their noise. At 7B, where the three learning arms sit within a few points of each other and the whole effect is 1.06–4.26 points, an analytic max-of- k argument can be made to cover most of the effect.

That argument assumes the arms are noisy *independently*. They are not: every arm branches from the same warm start, sees the same items, and fails on largely the same items, so the per-item correlation between arms is high and the variance of their maximum is far below the independent case. The question is empirical, and the per-item vectors settle it.

Table 12 reports two measurements. The first resamples items with replacement and computes Δ twice on each resample—once recomputing the arg max, once with the arm frozen at the full-sample arg max. The difference is exactly the bias that selecting and scoring on the same data introduces. It never exceeds 0.472 points in any group and averages -0.091 points.

The second is a cross-fitted estimator that removes the bias by construction: the control arm is selected on a random half of the test set and scored on the other half, averaged over $K = 2,000$ splits. Cross-fitted Δ is negative in **17 of 17 groups**, with a mean of -4.18 points against the naive -4.44 . Selection accounts for a quarter of a point of the four-point effect.

The reason is visible in Table 11: in the 11 groups at 0.5B and 1B the arg max is a *deterministic* arm with zero sampling variance that leads the nearest learning arm by at least 8.5 points, so no selection takes place at all. The winner’s-curse concern is real only in the six 7B groups, where three close learning arms compete—and there the measured bias is at most 0.47 points against effects of 1.06 to 4.26. The objection is sound in principle and empty in this data.

Table 13: Aggregator sensitivity, summary over the 17 groups. Only the pre-registered max over the full four-arm control group is negative everywhere; every weaker aggregator reports a positive mean effect.

control-group aggregator	mean Δ	range	$\Delta < 0$
max over $\mathcal{B}$ (pre-registered)	-4.44	-16.42 to -1.06	17/17
max over learning arms only	+5.66	-8.89 to +35.11	8/17
second best of $\mathcal{B}$	+6.43	-8.89 to +35.11	6/17
mean of $\mathcal{B}$	+7.97	-6.14 to +31.83	2/17
median of $\mathcal{B}$	+8.61	-5.80 to +40.07	2/17
solver alone	+1.13	-16.42 to +16.31	11/17
RS-SFT alone	+6.50	-8.89 to +35.11	6/17
continued SFT alone	+9.19	-4.26 to +45.04	8/17
warm start W_m alone	+15.07	+0.71 to +48.58	0/17

8.5 How much of the result is the solver?

The pre-registered rule aggregates the control group with a max. A reader is entitled to ask what any weaker aggregation would have said, and what the result owes to the one arm we built ourselves.

Table 13 answers both. Only the pre-registered max over the full four-arm group is negative everywhere. Against the mean of the control group the study reports +7.97 points; against the warm start alone, the comparison most common in the literature, it reports +15.07 points and is positive in every single group. That spread—from +15.07 to -4.44 on identical runs—is the paper’s central claim restated as a sensitivity analysis.

Removing the solver entirely is the most informative row. Against the maximum over the three learning arms, Δ averages +5.66 points but remains *negative in 8 of 17 groups*: both Config 2 $m=128$ groups and all six 7B groups. In other words, in nearly half the study the verdict does not depend on the solver at all; continued SFT alone is sufficient to overturn it. The solver is load-bearing in the 9 remaining groups, all at 0.5B and 1B. Per-group values are in Appendix A.

8.6 The two arms are not solving the same problems

The discordant counts in Table 11 carry information the Δ summary discards. In Config 1 at $m=2$, GRPO is right on 38 items the solver misses while the solver is right on 44 items GRPO misses: 82 of 282 items are decided differently by the two arms. Their union reaches 78.01% against the better arm’s 64.54%.

Across all 17 groups the union of GRPO and the best control arm exceeds the better of the two by 9.02 points on average (range 1.77 to 16.57), and the union of all five arms averages 75.37% against a best-single-arm average of 60.67%. This is not a deployable method—selecting the right arm per item requires the answer—but it bounds what the framing costs. Asking which of a policy-gradient arm and a symbolic arm is stronger, and reporting the winner, discards the fact that on these tasks they are substantially complementary. A control-group protocol that reports the max is still the right instrument for validating a claimed gain; it is not, on this evidence, the right instrument for deciding what to build.

9 Discussion

What generalizes. The specific numbers do not transfer beyond these two tasks; the design question does, and our scale sweep sharpens what that question is. The naive reading of the small-backbone result—“add a symbolic solver to your control group”—is already refuted by our own 7B data, where the solver is 15 points behind and contributes nothing. What survives both regimes is weaker and more portable: the control

group must be broad enough that *some* arm in it is competitive with RL under the conditions being reported, and it must be aggregated with a max so that the answer cannot be improved by choosing the opponent. Section 8 quantifies the cost of neglecting this on our own runs: the same 17 groups report +15.07 points against the warm start alone, +7.97 against the mean of the control group, and −4.44 against its maximum.

What the recommendation is not. A reasonable objection is that “aggregate with a max over a broad control group” is unfalsifiable as guidance, since a sufficiently determined critic can always append one more arm until any result is negative. We agree that the prescription needs a boundary, and we do not think this paper establishes where it lies. Two properties distinguish our use of it from an open-ended ratchet. The arm set was fixed by pre-registration before the confirmatory task ran, so it could not grow in response to the result; and every arm receives the same demonstration budget, so no arm can be strengthened by spending resources the proposed method was denied. A protocol that fixes its arm set in advance and equalizes budgets is a test. One that adds arms after seeing the outcome is not. We make no claim to have found the right stopping rule for how broad a control group must be.

Scope of the claim. Two tasks, three backbones spanning 14×, one RL algorithm (GRPO), one task family (program-annotated financial QA), and deliberately scarce demonstration budgets. Three narrower scope conditions deserve to be stated plainly rather than inferred:

One algorithm, not a family. Every RL number in this paper is GRPO with one implementation. We have no evidence about PPO, RLOO, or the several GRPO variants that correct its known estimator biases [8]. Where this paper says “RLVR” as shorthand, the evidence underneath it is GRPO.

One parameterization. Every arm—including every baseline—trains a LoRA adapter at $r=16$, $\alpha=32$, learning rate 10^{-5} . LoRA is applied symmetrically, so it is not an unfair handicap on the RL arm specifically, but it is a joint restriction on all arms and we cannot say whether full-parameter training would change the ordering. It was not listed as a threat in the conference version; it should have been.

The clean filter is our own, and it cannot manufacture the result. An item enters the study only if its gold program executes to its gold answer under our verifier (282/288 of the CodeTAT-QA test set, 664/795 of CodeFinQA). The same verifier defines the training reward and grades the evaluation, which is internally consistent but invites the charge that we selected the items on which our conclusion holds. It cannot work that way. The filter is applied identically to all five arms, so scoring every discarded item as wrong for everyone multiplies each arm’s accuracy by the same constant—0.979 for CodeTAT-QA, 0.835 for CodeFinQA—and a positive rescaling cannot change the sign of a difference. We verified this directly: recomputing all 17 groups on the unfiltered test sets leaves **17 of 17** signs unchanged, with Δ moving from −4.44 to −4.00 on average. No re-run was required; the objection fails on arithmetic.

External calibration. What the filter does affect is comparability of absolute numbers, so we place ours against the published BizBench baselines, which use the same 1%-relative-tolerance metric. On CodeTAT-QA our 7B arms reach 79–85%, above the reported 75.0% for Mistral-7B and 70.9% for CodeLlama-7B and in the band of CodeLlama-34B (81.4%) and Mixtral-8x7B (83.9%); our 0.5B GRPO reaches 63.12%, near Llama-2-13B (65.1%). On CodeFinQA our 1B arms reach up to 30.57%, comparable to MPT-30B (31.0%) and Llama-2-13B (33.4%). These are not weak systems for their size. The solver’s 64.54% sits between MPT-30B (64.8%) and Llama-2-13B (65.1%), which is itself a quantification of this paper’s point: 440 lines of rules match a 13–30B model’s few-shot performance on this task. One caveat about the comparison itself: our splits are the public BizBench release (kensho/bizbench), whose CodeFinQA and CodeTAT-QA test sets hold 795 and 288 items and whose training sets hold 4,668 and 2,856—all four matching our copies exactly—whereas the BizBench paper’s own table reports test sets of 844 and 392. We cannot reconcile that difference from the outside, so the published accuracies above should be read as the right order of magnitude for each model scale rather than as strictly paired comparisons.

Threats to validity. *The RL arm’s budget is the largest threat, and it is not resolved.* Section 7 reports that no GRPO run completed one epoch over its prompt pool (0.10–0.79), and that at 7B only 50 to 82 prompts per run produced a non-zero group advantage. Both numbers point the same direction: our GRPO arm is small by contemporary standards, and most visibly so at the scale where this paper’s generalization claim lives. The pre-registration fixed equal wall clock, which at a fixed per-token cost buys an order of magnitude less data at 7B than at 0.5B. We believe the 0.5B and 1B result is robust—there GRPO learns vigorously, beating every learning arm by up to 35 points, and still loses to the solver. We do not believe this study settles what a well-resourced GRPO run does at 7B, and readers should not take it to.

Solver strength as a free parameter. At 0.5B and 1B the solver sets the bar. Section 7.8 gives its development protocol, its dev trajectory, its 7.46-point dev-to-test penalty and its engineering cost, and the logs are published. The 7B result also removes this objection structurally: there the maximum is attained by continued SFT, an arm we did not design. Section 8 shows the verdict is negative in 8 of 17 groups with the solver deleted from the control group entirely.

The operating points were chosen inside the scarce regime. This is true and it is by design: the pre-registered bands select the two smallest budgets reaching stated warm-start accuracy windows, which places both operating points well left of the $m \approx 1024$ crossing where learning arms overtake the solver. The claim is therefore explicitly confined to the demonstration-scarce regime. What we cannot offer is a confirmatory group past the crossing point: no frozen-test evaluation was run at $m \geq 256$ in Config 1, so the statement that the effect reverses there rests on dev dose curves rather than on the decision rule.

Statistical power. Under Holm correction across the family of 17, two groups are significant. The per-group minimum detectable effect is a median of 6.93 points corrected, against measured effects with a median of 2.84. The quantitative claim rests on six independent operating points agreeing in sign ($p = 0.031$), not on per-group significance, and not on the count of 17.

Dev-selection bias is three observations, not a measurement. We report three instances where Δ is non-negative on dev and negative on the frozen test. Each is real and the $\beta=0$ case is instructive, but three instances identified after the fact do not establish the stated mechanism—that policy-gradient trajectories fluctuate more on dev than fine-tuning trajectories and so gain more from checkpoint selection. Measuring that would require per-arm dev-checkpoint variance across all groups, which our stored artifacts support only at 7B. We therefore state it as three observations consistent with a mechanism rather than as a measurement of one, and a reader should weigh contribution (iii) accordingly: what is established is the practical point that a frozen test set changes three verdicts, not the explanation for why.

Low absolute accuracy in Config 2. The solver reaches only $\sim 32\%$ there, so all arms operate in a low band. This is why the study spans a two-fold difference in solver strength and a $14\times$ range in backbone; the sign is unchanged across all of it.

One exclusion. All 85 arm evaluations of the 17 reported groups are present. The single excluded group is the collapsed run of Section 7, whose Δ (-19.73) is more negative than any we report, so excluding it weakens rather than strengthens the result. The pre-registration did not specify an exclusion rule for training collapse; we adopted one after observing it, and we report the result both ways.

What would change our mind, and what did. A null result should say what would refute it. We named three things. A GRPO run at this scale with an order of magnitude more prompt coverage that beats $\max \mathcal{B}$ at any operating point. A confirmatory group past the $m \approx 1024$ crossing in which GRPO leads a control group including the solver. Or a task family in the same program-annotated setting where no deterministic arm is viable and GRPO still fails to beat continued SFT, which would indicate our result is about the solver rather than about the control group.

The first one fired. At ten times the prompt coverage, the smallest-margin 7B cell reverses from $\Delta = -1.06$ to $+6.74$ and $+6.38$ across two training seeds, clearing our own Go threshold with McNemar $p =$

0.0013 and 0.0029 (Appendix G). We report it in the abstract, in Section 6, in Section 7 and here, because a paper that specifies its own refutation conditions and then meets one of them owes the reader the result in the same places it made the claim.

What survives is bounded rather than broken. The reversal is a budget effect, not a defect effect: a matched-budget run that raises informative prompts from 82 to 120 lands at +1.06 with an interval of $[-2.13, +4.26]$, statistically where the registered -1.06 already was. So under the pre-registered protocol—equal wall clock for every arm—the finding holds at 7B; given ten times that budget it does not. Config 3 therefore marks where this study’s claim ends, and the quantitative core is the 11 groups at 0.5B and 1B. The other two tests we have not attempted.

On reporting negative results. The pre-registered protocol accepted both outcomes in advance, including explicit language forbidding us from weakening the solver if GRPO won. We mention this because a null result is only informative if the analysis could have come out otherwise. In the same spirit, this extended version reports the two findings that most damage our own case—two of seventeen groups significant under our own correction, and an RL arm that never completed an epoch—in the main text rather than in an appendix.

10 Conclusion

We pre-registered a decision rule that compares GRPO against the maximum of a four-arm control group, one arm of which does not learn, and executed it on two tasks and three backbones spanning $14\times$ in scale. All 17 frozen-test effect sizes are negative, from -1.06 to -16.42 points. At 0.5B, the same runs judged against a learning-only control group would have been reported as substantial successes—GRPO beats every learning baseline there by up to 35 points.

The mechanism is not one strong baseline but a moving one. At 0.5B and 1B a deterministic solver attains the maximum in all 11 groups, because it is flat in the demonstration budget while every learning arm scales with it. At 7B the solver falls 15 points behind and drops out of contention—exactly as the “gains emerge with capability” account predicts—and Δ remains negative because continued SFT takes its place. Neither ablations (the solver is 100% rule-derived) nor audits (zero constant-program shortcuts in 57,517 positive samples) nor a hyperparameter sweep provides an escape; the sweep’s one positive Δ appears on dev and reverses on the frozen test.

This extended version adds the audit of our own instrument. The max in the decision rule does not manufacture the result: a cross-fitted estimator that selects the control arm on one half of the test set and scores it on the other is negative in 17 of 17 groups. The sampling arm usually proposed as the missing baseline is not deployable here, and even its oracle bound loses to the solver at three of four operating points. Against that, only two groups survive Holm correction, no GRPO run completed an epoch, and—most consequentially—the 7B configuration reverses at ten times the budget, so it bounds the claim rather than extending it. Within the registered protocol the negative sign is robust; the per-group effect sizes are not; and the scale evidence is a budget boundary.

What generalizes is not the sign but the dependence. This study can now put three numbers on how much a reported RLVR effect moves without any change to the policy being evaluated: 19.5 points for which control arm is nominated, 7.8 for an order of magnitude of compute, and 3.0 for the inference engine alone. An RLVR result reported without its comparison, its budget and its decoding stack attached is not yet a result.

The practical recommendation is therefore not “add a solver.” It is: aggregate your control group with a maximum rather than nominating one opponent, populate it broadly enough that the aggregation can bind,

fix the arm set before you see the result, and reserve a test set that selection never touches. Which arm ends up binding will depend on the regime—a solver in ours at small scale, plain continued SFT at 7B—and that is precisely why the choice cannot be made in advance by the party reporting the gain.

References

- [1] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo, “DeepSeekMath: Pushing the Limits of Mathematical Reasoning in Open Language Models,” *arXiv preprint arXiv:2402.03300*, 2024.
- [2] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al., “DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning,” *arXiv preprint arXiv:2501.12948*, 2025.
- [3] Zheng Yuan, Hongyi Yuan, Chengpeng Li, Guanting Dong, Keming Lu, Chuanqi Tan, Chang Zhou, and Jingren Zhou, “Scaling Relationship on Learning Mathematical Reasoning with Large Language Models,” *arXiv preprint arXiv:2308.01825*, 2023.
- [4] Avi Singh, John D. Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, Xavier Garcia, Peter J. Liu, James Harrison, Jaehoon Lee, Kelvin Xu, Aaron Parisi, Abhishek Kumar, Alex Alemi, Alex Rizkowsky, Azade Nova, Ben Adlam, Bernd Bohnet, Gamaleldin Elsayed, Hanie Sedghi, Igor Mordatch, Noah Fiedel, Peter Hawkins, Rosanne Liu, Simon Kornblith, Noah Fiedel, Orhan Firat, Matthew D. Hoffman, Aleksandra Faust, and Sanjiv Kumar, “Beyond Human Data: Scaling Self-Training for Problem-Solving with Language Models,” *arXiv preprint arXiv:2312.06585*, 2023.
- [5] Zhiyu Chen, Wenhui Chen, Charese Smiley, Sameena Shah, Iana Borova, Dylan Langdon, Reema Moussa, Matt Beane, Ting-Hao Huang, Bryan Routledge, and William Yang Wang, “FinQA: A Dataset of Numerical Reasoning over Financial Data,” in *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2021.
- [6] Fengbin Zhu, Wenqiang Lei, Youcheng Huang, Chao Wang, Shuo Zhang, Jiancheng Lv, Fuli Feng, and Tat-Seng Chua, “TAT-QA: A Question Answering Benchmark on a Hybrid of Tabular and Textual Content in Finance,” in *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021.
- [7] Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman, “STaR: Bootstrapping Reasoning With Reasoning,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [8] Zichen Liu et al., “Understanding R1-Zero-Like Training: A Critical Perspective,” *arXiv preprint arXiv:2503.20783*, 2025.
- [9] Weihao Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He, “SimpleRL-Zoo: Investigating and Taming Zero Reinforcement Learning for Open Base Models in the Wild,” in *Conf. Language Modeling (COLM)*, 2025.
- [10] Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang, “Does Reinforcement Learning Really Incentivize Reasoning Capacity in LLMs Beyond the Base Model?,” *arXiv preprint arXiv:2504.13837*, 2025.
- [11] Andreas Hochlehnert, Hardik Bhatnagar, Vishaal Udandara, Samuel Albanie, Ameya Prabhu, and Matthias Bethge, “A Sober Look at Progress in Language Model Reasoning: Pitfalls and Paths to Reproducibility,” *arXiv preprint arXiv:2504.07086*, 2025.

- [12] Maurizio Ferrari Dacrema, Paolo Cremonesi, and Dietmar Jannach, “Are We Really Making Much Progress? A Worrying Analysis of Recent Neural Recommendation Approaches,” in *Proc. 13th ACM Conf. Recommender Systems (RecSys)*, 2019, pp. 101–109.
- [13] Kevin Musgrave, Serge Belongie, and Ser-Nam Lim, “A Metric Learning Reality Check,” in *Proc. European Conf. Computer Vision (ECCV)*, 2020, pp. 681–699.
- [14] Jimmy Lin, “The Neural Hype and Comparisons Against Weak Baselines,” *ACM SIGIR Forum*, vol. 52, no. 2, pp. 40–51, 2018.
- [15] Peter Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, and David Meger, “Deep Reinforcement Learning That Matters,” in *Proc. AAAI Conf. Artificial Intelligence*, 2018.
- [16] Rishabh Agarwal, Max Schwarzer, Pablo Samuel Castro, Aaron C. Courville, and Marc G. Bellemare, “Deep Reinforcement Learning at the Edge of the Statistical Precipice,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- [17] Cynthia Dwork, Vitaly Feldman, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Aaron Roth, “The reusable holdout: Preserving validity in adaptive data analysis,” *Science*, vol. 349, no. 6248, pp. 636–638, 2015.
- [18] Benjamin Recht, Rebecca Roelofs, Ludwig Schmidt, and Vaishaal Shankar, “Do ImageNet Classifiers Generalize to ImageNet?,” in *Proc. 36th Int. Conf. Machine Learning (ICML)*, 2019, pp. 5389–5400.
- [19] Sture Holm, “A Simple Sequentially Rejective Multiple Test Procedure,” *Scandinavian Journal of Statistics*, vol. 6, no. 2, pp. 65–70, 1979.
- [20] Dallas Card, Peter Henderson, Urvashi Khandelwal, Robin Jia, Kyle Mahowald, and Dan Jurafsky, “With Little Power Comes Great Responsibility,” in *Proc. Conf. Empirical Methods in Natural Language Processing (EMNLP)*, 2020, pp. 9263–9274.
- [21] Wenhui Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen, “Program of Thoughts Prompting: Disentangling Computation from Reasoning for Numerical Reasoning Tasks,” *arXiv preprint arXiv:2211.12588*, 2022.
- [22] Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig, “PAL: Program-aided Language Models,” in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023.
- [23] Zhiyu Chen, Shiyang Li, Charese Smiley, Zhiqiang Ma, Sameena Shah, and William Yang Wang, “ConvFinQA: Exploring the Chain of Numerical Reasoning in Conversational Finance Question Answering,” in *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2022.
- [24] Rik Koncel-Kedziorski, Michael Krumdick, Viet Lai, Varshini Reddy, Charles Lovering, and Chris Tanner, “BizBench: A Quantitative Reasoning Benchmark for Business and Finance,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- [25] Chen Liang, Jonathan Berant, Quoc Le, Kenneth D. Forbus, and Ni Lao, “Neural Symbolic Machines: Learning Semantic Parsers on Freebase with Weak Supervision,” in *Proc. 55th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2017, pp. 23–33.

- [26] Chen Liang, Mohammad Norouzi, Jonathan Berant, Quoc Le, and Ni Lao, “Memory Augmented Policy Optimization for Program Synthesis and Semantic Parsing,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.
- [27] Panupong Pasupat and Percy Liang, “Compositional Semantic Parsing on Semi-Structured Tables,” in *Proc. 53rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2015, pp. 1470–1480.
- [28] Zhoujun Cheng et al., “Binding Language Models in Symbolic Languages,” in *Int. Conf. Learning Representations (ICLR)*, 2023.
- [29] Zhaowei Liu, Xin Guo, Zhi Yang, Fangqi Lou, Lingfeng Zeng, Jinyi Niu, Mengping Li, Qi Qi, Zhiqiang Liu, Yiyang Han, Dongpo Cheng, Ronghao Chen, Huacan Wang, Xingdong Feng, Huixia Judy Wang, Chengchun Shi, and Liwen Zhang, “Fin-R1: A Large Language Model for Financial Reasoning through Reinforcement Learning,” *arXiv preprint arXiv:2503.16252*, 2025.
- [30] Lingfei Qian, Yan Wang, Xueqing Peng, Weipeng Zhou, Yi Han, Yupeng Cao, Yilun Zhao, Guojun Xiong, Jimin Huang, Qianqian Xie, and Jian-Yun Nie, “Fino1: On the Transferability of Reasoning-Enhanced LLMs and Reinforcement Learning to Finance,” *arXiv preprint arXiv:2502.08127*, 2025.
- [31] Jie Zhu, Qian Chen, Huaixia Dou, Junhui Li, Lifan Guo, Feng Chen, and Chi Zhang, “DianJin-R1: Evaluating and Enhancing Financial Reasoning in Large Language Models,” *arXiv preprint arXiv:2504.15716*, 2025.

A Every frozen-test number

Table 14 gives all 17 groups with no averaging over seeds: five arms, the arm attaining $\max \mathcal{B}$, and Δ . Table 1 in the main text is this table averaged within cells.

Table 15 gives the per-group aggregator sensitivity summarized in Table 13. Reading down the *learn* column identifies the 8 groups in which the verdict does not depend on the solver at all: Config 2 at $m=128$ (both surviving seeds) and all six Config 3 groups. Reading down the *solver* column identifies the 11 groups in which the solver is the binding arm.

B Complete solver ablation

Variant **A** uses only the hand-written rule path, **B** only retrieval, and **C** is the deployed configuration, which tries the rule path first and falls back to retrieval. If the solver’s accuracy came from copying gold programs out of its retrieval library it would be a memorization artifact rather than a baseline, so the quantity of interest is $C - A$: what retrieval adds on top of the rules.

Variant **B** returns 0.00% at every one of Config 1’s ten budgets and at every Config 2 budget up to and including $m=256$: it never produces an accepted program. Across all ten Config 1 budgets and three seeds—thirty measurements— A , C and their difference are constant at 72.00%, 72.00% and +0.00 points.

Config 2 is the only place retrieval ever fires, and it does so only at $m=1024$ (1.00%) and at the full pool $m=3635$ (2.00%). Both lie far above the two operating points from which reported groups are drawn ($m_L=128$, $m_H=256$); at those two the deployed variant differs from the rule-only path by -0.67 and -0.33 points, which is one to two items on the 200-item dev subset.

With retrieval never firing at any operating budget, the demonstration budget has no channel through which to affect the solver at all. That is the mechanism behind the flat dose curve in the main paper: the solver is not merely insensitive to m in practice, it is structurally unable to use m .

Table 14: Every frozen-test number in the study: all 17 groups, all five arms, no averaging over seeds. Config 2 $m=128$ seed 1 is the excluded collapsed run.

m	seed	solver	W_m	c-SFT	RS-SFT	GRPO	arg max	Δ
<i>Config 1 (CodeTAT-QA $\times 0.5B$, $n=282$)</i>								
2	0	64.54	15.60	24.47	27.66	62.41	solver	-2.13
2	1	64.54	14.54	18.09	28.01	63.12	solver	-1.42
2	2	64.54	26.60	20.57	30.50	60.64	solver	-3.90
64	0	64.54	29.43	43.62	51.06	58.51	solver	-6.03
64	1	64.54	27.66	45.39	56.03	60.99	solver	-3.55
64	2	64.54	32.62	45.74	51.42	61.70	solver	-2.84
<i>Config 2 (CodeFinQA $\times 1B$, $n=664$)</i>								
128	0	32.38	12.05	20.18	21.84	17.02	solver	-15.36
128	2	32.53	12.65	18.83	25.00	16.11	solver	-16.42
256	0	32.38	25.60	26.36	28.92	29.97	solver	-2.41
256	1	32.83	25.75	27.11	25.90	28.16	solver	-4.67
256	2	32.08	28.16	27.41	30.27	30.57	solver	-1.51
<i>Config 3 (CodeTAT-QA $\times 7B$, $n=282$)</i>								
2	0	64.54	71.28	79.79	75.18	76.60	c-SFT	-3.19
2	1	64.54	69.86	77.66	76.95	74.82	c-SFT	-2.84
2	2	64.54	70.57	76.95	74.82	75.18	c-SFT	-1.77
64	0	64.54	79.08	81.91	80.50	79.79	c-SFT	-2.13
64	1	64.54	78.72	80.50	80.14	79.43	c-SFT	-1.06
64	2	64.54	79.43	85.11	81.21	80.85	c-SFT	-4.26

C Non-learning alternatives: prompting and sampling

C.1 Program-of-Thought / PAL prompting

The first alternative non-learning baseline is to prompt the base model directly to emit a program, using the same m demonstrations as few-shot examples and taking no gradient steps.

Frozen-test accuracy is 9.93–16.67% (mean 13.30%) in Config 1 and 0.75–1.36% (mean 1.00%) in Config 2. This is the weakest arm in the entire study, far below both the solver and GRPO, which is what establishes that the solver’s strength is specific symbolic competence rather than a free advantage of skipping training.

Two features of the table need explaining. First, the Config 2 rows at $m=128$ and $m=256$ are identical. This is not a transcription error: the prompt carries at most four few-shot examples, and because the demonstration subsets are nested prefixes, both budgets supply the same four. Second, the Config 1 $m=64$ seed-0 dev evaluation ran on a 20-item subset rather than the 200-item convention used everywhere else; its frozen-test evaluation used the full 282 items, as did all others.

C.2 Best-of- n and majority voting

The second alternative is inference-time sampling from the warm start, discussed in Section 7.7. Because our verifier requires the gold answer, verifier-filtered best-of- n coincides with pass@ n and is an oracle bound rather than a deployable arm; majority voting is the deployable relative.

At the four pre-registered operating budgets the deployable arm is below the solver everywhere, and the oracle bound is below the solver at three of the four. Both sampling arms overtake the solver from $m=256$ (Config 1) and $m=256$ (Config 2 oracle) onward, which is the same crossing the dose curves show for every

Table 15: Aggregator sensitivity, per group. Same runs throughout; only the way the control group is reduced to one number changes.

m	seed	max $\mathcal{B}$	max learn	c-SFT	RS-SFT	W_m	solver	mean	2nd
<i>Config 1 (CodeTAT-QA $\times 0.5B$, $n=282$)</i>									
2	0	-2.13	+34.75	+37.94	+34.75	+46.81	-2.13	+29.34	+34.75
2	1	-1.42	+35.11	+45.04	+35.11	+48.58	-1.42	+31.83	+35.11
2	2	-3.90	+30.14	+40.07	+30.14	+34.04	-3.90	+25.09	+30.14
64	0	-6.03	+7.45	+14.89	+7.45	+29.08	-6.03	+11.35	+7.45
64	1	-3.55	+4.96	+15.60	+4.96	+33.33	-3.55	+12.59	+4.96
64	2	-2.84	+10.28	+15.96	+10.28	+29.08	-2.84	+13.12	+10.28
<i>Config 2 (CodeFinQA $\times 1B$, $n=664$)</i>									
128	0	-15.36	-4.82	-3.16	-4.82	+4.97	-15.36	-4.59	-4.82
128	2	-16.42	-8.89	-2.71	-8.89	+3.46	-16.42	-6.14	-8.89
256	0	-2.41	+1.05	+3.61	+1.05	+4.37	-2.41	+1.66	+1.05
256	1	-4.67	+1.05	+1.05	+2.26	+2.41	-4.67	+0.26	+1.05
256	2	-1.51	+0.30	+3.16	+0.30	+2.41	-1.51	+1.09	+0.30
<i>Config 3 (CodeTAT-QA $\times 7B$, $n=282$)</i>									
2	0	-3.19	-3.19	-3.19	+1.42	+5.32	+12.06	+3.90	+1.42
2	1	-2.84	-2.84	-2.84	-2.13	+4.96	+10.28	+2.57	-2.13
2	2	-1.77	-1.77	-1.77	+0.35	+4.61	+10.64	+3.46	+0.35
64	0	-2.13	-2.13	-2.13	-0.71	+0.71	+15.25	+3.28	-0.71
64	1	-1.06	-1.06	-1.06	-0.71	+0.71	+14.89	+3.46	-0.71
64	2	-4.26	-4.26	-4.26	-0.35	+1.42	+16.31	+3.28	-0.35

learning arm: sampling more from a better-fit warm start is a learning arm in disguise, and it scales with the demonstration budget for the same reason the others do.

These are dev measurements on the 200-item selection subset and were not part of the pre-registered control group $\mathcal{B}$; they do not enter Δ .

D Verifier-hacking audit

If GRPO were exploiting the reward checker rather than solving the task, its accuracy would be an artifact and the comparison meaningless. We audit the reward-dump records of the 14 GRPO runs whose dumps were retained, spanning all three backbones. H1 flags a constant program that performs no computation; H2 flags the gold value appearing in the emitted code as a literal; H3 is the share of positive samples falling under the single most frequent code skeleton; H4 flags no-op arithmetic such as $x \times 1$ or $x + 0$, which would slip past the anti-constant guard.

Constant-program shortcuts occur **zero** times in 57,517 positive-reward samples: the verifier’s anti-constant guard holds. No-op arithmetic appears 27 times, 0.05% of positive samples. Template collapse is absent—the most frequent skeleton accounts for at most 2.1% of positive samples, with skeleton Gini coefficients between 0.06 and 0.45, indicating a broadly distributed solution space rather than a single exploited pattern.

H2 fires on 0.19% of positive samples. This is expected rather than suspicious on these tasks: gold answers are computed from numbers that appear in the context, so a correct program will sometimes contain a literal that coincides with the gold value.

The four unretained dumps are Config 3 runs. Their positive-reward rates and group statistics are nonetheless available from the run metadata and appear in Table 9, so no run is unaccounted for. The audit therefore covers 14 of 18 runs by dump inspection and 18 of 18 by recorded aggregate.

Table 16: Complete solver ablation. **A** uses only the hand-written rule path, **B** only retrieval, **C** is the deployed configuration (rules first, retrieval as fallback). Dev accuracy (%), mean over seeds, at every budget of each configuration. Retrieval never fires at any budget used for a reported group.

cfg	m	A: rules	B: retr.	C: deployed	$C - A$
1	0	72.00	0.00	72.00	+0.00
	1	72.00	0.00	72.00	+0.00
	2	72.00	0.00	72.00	+0.00
	4	72.00	0.00	72.00	+0.00
	8	72.00	0.00	72.00	+0.00
	16	72.00	0.00	72.00	+0.00
	64	72.00	0.00	72.00	+0.00
	256	72.00	0.00	72.00	+0.00
	1024	72.00	0.00	72.00	+0.00
	2446	72.00	0.00	72.00	+0.00
2	0	34.50	0.00	34.50	+0.00
	16	35.67	0.00	35.50	−0.17
	64	35.83	0.00	35.17	−0.67
	128	36.00	0.00	35.33	−0.67
	256	35.67	0.00	35.33	−0.33
	1024	36.00	1.00	36.33	+0.33
	3635	36.00	2.00	38.00	+2.00

E Training stability across all 18 GRPO runs

One run marked [†]—Config 2, $m=128$, seed 1—consumed its full wall-clock and verifier budget without ever receiving a positive reward: the policy drifted into a region where it no longer emits an executable program and did not recover. Its recorded metadata confirms the mechanism exactly: of 880 prompt groups, 880 are all-zero and none is all-one (Table 9). It is the single group excluded from the main paper, on the grounds that this is an infrastructure or training failure rather than a property of the algorithm under test. We note in Section 9 that the pre-registration did not specify this exclusion rule in advance.

The collapse is confined to the weakest backbone at the scarcest budget. Every other run sustains a positive reward rate throughout, and the rate rises monotonically with backbone capability: 0.522–0.560 at 0.5B, 0.081–0.092 and 0.194–0.239 at 1B, and 0.772–0.846 at 7B. The excluded run’s Δ would have been −19.73 points, more negative than any value we report, so excluding it weakens rather than strengthens the paper’s conclusion.

Section 7 explains why the monotone rise in positive-reward rate must not be read as a monotone rise in training quality: at 7B the high rate is produced by all-correct groups, which carry zero advantage. Table 9 gives the decomposition for every run.

F Pre-registration, deviations and reproduction

F.1 What was fixed, and when

The protocol was not written in one sitting, and describing it as a single pre-registration would overstate it. The sequence was:

Note first that the configuration numbers in this paper are not chronological. Config 2 ran first; Configs 1 and 3 are the second task.

Table 17: Program-of-Thought / PAL prompting of the base model, with the same m demonstrations as few-shot examples and no gradient steps. Accuracy (%) per seed on both splits. This is the weakest arm in the entire study.

cfg	m	seed	dev	n	test	n
1	2	0	14.00	200	12.77	282
	2	1	24.50	200	13.12	282
	2	2	26.00	200	16.67	282
	64	0	10.00	20	9.93	282
	64	1	19.00	200	12.77	282
	64	2	21.50	200	14.54	282
2	128	0	0.50	200	0.90	664
	128	1	0.50	200	0.75	664
	128	2	2.00	200	1.36	664
	256	0	0.50	200	0.90	664
	256	1	0.50	200	0.75	664
	256	2	2.00	200	1.36	664

1. A single-seed *smoke test* was frozen on 2026-07-19 and run on CodeFinQA. Its decision quantity was GRPO minus matched SFT—one opponent, not a max—with a Go bar at five points and a No-Go floor at three. Its own text states that single-seed results are smoke evidence and not paper evidence, and no number from it appears here.
2. The protocol this paper reports was written on 2026-07-20 as a new question rather than a patch on the first. It introduced the five arms, the max aggregation over $\mathcal{B}$, the +3 Go criterion with its companion conditions, the operating-band rule for m_L and m_H , and the good-faith clause forbidding us from weakening the solver if GRPO won. It was written for CodeFinQA, which this paper calls **Config 2**.
3. Config 2 was executed.
4. An addendum was frozen on 2026-07-24, *before any CodeTAT-QA training or evaluation run*, instantiating CodeTAT-QA as the second task and carrying the Gate 0/1/2 criteria, the definition of Δ and the verdict vocabulary over verbatim.
5. Configs 1 and 3 (CodeTAT-QA at 0.5B and 7B) were executed under it.

The honest labelling that follows is the reverse of what the numbering suggests. **Config 2 is the exploratory one**: its protocol was written by authors who had already seen single-seed results on that same dataset. **Configs 1 and 3 are confirmatory**: the rule governing them was frozen before any of their data existed. The abstract’s framing—“pre-registered, before the second task was run”—refers to step 4 and is accurate, but a reader could take it to mean the whole study was registered before any data existed, and it was not. Note also that the quantitative core of this paper, the 11 groups at 0.5B and 1B, straddles the two: six of them are confirmatory and five are not.

One further qualification on the word “frozen.” The document written in step 2 carries a status line reading *not yet frozen*, and the addendum of step 4 refers to the operative protocol as that document *together with* a numbered entry in the deviation log. The frozen protocol is therefore a draft plus one logged amendment rather than a single signed artifact, and a reader reconstructing it must read both. We publish both.

Within the confirmatory scope, only “GRPO vs. the strongest non-RL arm” is a confirmatory comparison; the ablations, audits, sweep, dose curves and every analysis in Section 8 of this extended version are exploratory and are reported as such.

Table 18: Inference-time sampling from the warm start, $n=16$, dev. pass@16 is an oracle bound because the verifier needs the gold answer; maj@16 is the deployable arm. * marks the two pre-registered operating budgets. The solver column is its dev accuracy at the same budget.

m	seeds	pass@1	pass@16	maj@16	solver	oracle beats solver?
<i>Config 1 (CodeTAT-QA $\times$ 0.5B, $n=282$)</i>						
0	3	3.33	22.00	12.50	72.00	no
1	3	11.67	47.83	32.67	72.00	no
2*	3	9.83	48.83	35.33	72.00	no
4	1	11.50	54.00	39.00	72.00	no
8	1	16.00	55.50	41.50	72.00	no
16	3	13.33	54.50	39.67	72.00	no
64*	3	13.17	61.17	45.17	72.00	no
256	3	55.33	84.33	73.00	72.00	yes
1024	3	66.00	91.50	78.17	72.00	yes
2446	3	69.83	94.67	81.83	72.00	yes
<i>Config 2 (CodeFinQA $\times$ 1B, $n=664$)</i>						
0	1	0.50	1.00	0.50	34.50	no
16	1	0.50	18.00	5.00	35.50	no
64	1	0.50	14.50	4.00	35.17	no
128*	3	5.50	34.50	17.17	35.33	no
256*	3	16.00	51.67	30.67	35.33	yes
1024	1	24.50	65.00	46.00	36.33	yes
3635	1	35.00	71.50	50.00	38.00	yes

Table 19: Verifier-hacking audit over the 14 GRPO runs whose reward dumps were retained. H1: constant program; H2: gold value appearing as a literal; H3: share of positive samples under the single most frequent code skeleton; H4: no-op arithmetic. Counts are over positive-reward samples.

cfg	runs	samples	positive	H1	H2	H4	H3
1	6	87,872	47,616	0	90	4	$\leq 0.9\%$
2	6	37,933	6,030	0	14	23	$\leq 1.5\%$
3	2	4,672	3,871	0	7	0	$\leq 2.1\%$
total	14	130,477	57,517	0	111	27	$\leq 2.1\%$

One rule was adopted after the fact: the exclusion of a run whose positive reward rate is identically zero. It was not specified in advance. Its effect is reported in both directions in Appendix E.

F.2 Artifacts

The following are released with this extended version, at <https://github.com/YOUR-ACCOUNT/rlvr-control-group-audit>:

- **The pre-registration documents and their addendum**, verbatim, in the state they were frozen.
- **The deviation logs**, listing every procedural departure from the protocol and the direction in which it influences Δ . These are not curated: they include the incident in which an early verdict script silently omitted the solver from the max set and reported $\Delta = +3.00$ where the corrected value is -14.00 . That error is why the aggregation is computed over an explicit arm list and checked, and we regard publishing it as part of the claim.

Table 20: Positive-reward rate across all 18 GRPO runs, with optimizer steps and verifier calls. One run—the weakest backbone at the scarcest budget—ends at zero and is the single excluded group.

backbone	m	seed	steps	verifier calls	pos. rate
Qwen2.5-0.5B	2	0	242	15,488	0.5266
	2	1	280	17,920	0.5602
	2	2	258	16,512	0.5509
	64	0	190	12,160	0.5222
	64	1	197	12,608	0.5439
	64	2	206	13,184	0.5398
Llama-3.2-1B	128	0	98	6,272	0.0807
	128	1 [†]	110	7,040	0.0000
	128	2	102	6,528	0.0921
	256	0	111	7,104	0.1958
	256	1	111	7,104	0.1944
	256	2	123	7,872	0.2387
Qwen2.5-7B	2	0	38	2,432	0.8121
	2	1	36	2,304	0.7721
	2	2	45	2,880	0.8233
	64	0	35	2,240	0.8464
	64	1	39	2,496	0.8369
	64	2	35	2,240	0.8263

- **The solver development logs**, round by round, with the dev accuracy after each round and the failure mode each round targeted (Section 7.8).
- **The solver, the verifier and the training/evaluation harness.**
- **Per-item outcome vectors** for all five arms in all 17 groups: the inputs to every number in Section 8 and Appendix A.
- **The inference-engine equivalence gate** of Section F.3: its script, its raw result and its 18 per-item flip samples.
- **The analysis scripts.** Every table in this document that reports a statistic is emitted directly from the stored artifacts by `gen_extended_stats.py` or `gen_supp_tables.py` under `analysis/`; both carry self-checks that abort if a recomputed value disagrees with the value printed in the paper.

F.3 The inference-engine equivalence gate

One entry in the deviation log bears directly on whether a 3-point Go threshold is measurable at all, so we state it here rather than leaving it to a reader of the logs.

Early in the study we considered replacing HuggingFace `generate` with vLLM, which is roughly $60\times$ faster on this workload. Before adopting it we ran an equivalence gate: the same checkpoint, the same 200 dev items, a prompt fingerprint verified identical on both sides, and—by design—generation separated from verification, so that the two engines only produced text and a single shared verifier scored it. Anything else would have folded environment differences into the measurement.

HuggingFace generate accuracy	41.00%
vLLM accuracy	38.00%
difference	−3.00 pp
per-item reward agreement	91.00% (18/200 flipped)
generated text identical	38.50%
HF-only correct / vLLM-only correct	12 / 6
McNemar χ^2	1.39 (n.s.)

The engine alone moves accuracy by 3.00 points, which is exactly this study’s pre-registered Go threshold, and it does so in no systematic direction. A constant offset could have been corrected; non-systematic noise of the same magnitude as the effect under test cannot. The mechanism is not floating-point jitter at the last digit but path divergence: a 0.5B policy reading a ~ 3000 -token prompt in bf16 has very small logit gaps, so a single `argmax` flip early in decoding sends the emitted program down a different branch. Only 38.5% of completions are textually identical, which shows the divergence is the normal case rather than an occasional accident.

We therefore rejected the substitution and used HuggingFace `generate` throughout. We report this for two reasons. First, it is the reason a reader should believe the 3-point threshold means anything in this study: we checked that our own instrument does not move by that much for reasons unrelated to the treatment. Second, it generalizes the paper’s argument one level down. The headline result is that a reported RLVR effect depends on which control arm you compare against, to the tune of 19 points; this gate shows that it also depends on which inference engine you decode with, to the tune of 3—and that an author who swapped engines mid-study would see neither dependence in the reported numbers.

F.4 Reproducing the numbers

With the artifact bundle in place:

```
# results_backup/ -> the fact table
python analysis/gen_data_facts.py
# ablation, PoT, audit, stability
python analysis/gen_supp_tables.py
# per-group stats, aggregators, cross-fitting,
# coverage, zero-advantage
python analysis/gen_extended_stats.py
```

Randomized quantities are seeded: the bootstrap uses $B = 10,000$ resamples and the cross-fitted estimator $K = 2,000$ splits, both under a fixed seed, so the reported intervals are reproducible to the digit.

F.5 Hardware and software

One NVIDIA RTX A6000 (48 GB) per run; PyTorch 2.10, transformers 4.57, peft 0.19, trl 1.8. Every arm trains a LoRA adapter ($r=16$, $\alpha=32$) at learning rate 10^{-5} ; GRPO uses 8 generations per prompt, $\beta=0.04$, temperature 1.0, top- p 0.95 and effective batch 64, reduced to per-device 4 with accumulation 16 at 7B. The three arms that keep training are each capped at 5700 s of wall clock. The solver runs on CPU in seconds.

Table 21: Follow-up runs on the Config 3 cell with the smallest registered margin ($m=64$, seed 1), where $\max \mathcal{B}$ is continued SFT at 80.50%. All three are exploratory and do not enter the pre-registered decision quantity. Coverage is prompt groups divided by the 2382-item pool.

run	steps	wall (s)	coverage	zero-var	informative	test acc	Δ
registered	39	5750	0.13	73.7%	82	79.43	-1.06
10 $\times$ coverage, seed A	414	57667	1.39	86.0%	462	87.23	+ 6.74
10 $\times$ coverage, seed B	415	57694	1.39	86.1%	461	86.88	+ 6.38
difficulty-filtered, matched budget	33	5753	0.11	54.5%	120	81.56	+ 1.06

G Follow-up: where the 7B result ends

Section 9 names three things that would change our mind. The first was “a GRPO run at this scale with an order of magnitude more prompt coverage, reaching a plateau in its own reward curve, that beats $\max \mathcal{B}$ at any operating point.” We ran it. It does.

This appendix reports that result and a second run that isolates why. Both are exploratory: they use the frozen test set at an operating point the protocol already spent, they were specified after the 17 registered groups were complete, and they do not enter Δ or any count in the main text. We report them because we said we would, and because the answer bounds the paper’s claim in a way a reader is entitled to know.

G.1 Design

We retested the single cell with the *smallest* registered margin: Config 3 at $m=64$, seed 1, where $\max \mathcal{B}$ is continued SFT at 80.50% and the registered GRPO run scored 79.43% for $\Delta = -1.06$. This is the cell most favourable to GRPO and therefore the sharpest test: if more compute cannot flip it here, it will not flip the others.

Three runs, all branching from the same warm start W_{64} on the same prompt pool, differing only in budget and in which prompts they see:

- **Two runs at ten times the wall clock** (57,600 s against 5700 s), with two different training seeds and the demonstration seed held fixed, so W_m and the pool are identical to the registered run.
- **One run at the registered budget with a difficulty-filtered pool.** Section 7 shows that 73.7% of this cell’s prompt groups carry zero advantage. We scored the pool once with W_{64} under greedy decoding, kept only the 440 of 2382 prompts it answered incorrectly, and trained on that subset under the *same* 5700 s cap. This asks whether the registered result is an artifact of wasted gradient signal rather than of budget.

G.2 What happened

At ten times the budget, GRPO wins, decisively and reproducibly. Both seeds reach 87.23% and 86.88% against continued SFT’s 80.50%: $\Delta = +6.74$ and $+6.38$, McNemar $p = 0.0013$ and 0.0029 , bootstrap intervals $[+2.84, +10.64]$ and $[+2.48, +10.28]$, and a bootstrap probability of clearing the +3 Go threshold of 0.97 and 0.95. By the paper’s own pre-registered criterion this is a Go. Coverage rose from 0.13 epochs to 1.39.

At the registered budget, fixing the degeneracy changes nothing measurable. The difficulty filter worked as intended: zero-advantage groups fell from 73.7% to 54.5% and informative prompts rose from

82 to 120, while the run consumed *less* on every compute axis than the registered one (33 steps against 39, 2112 verifier calls against 2496, 5753 s against 5750 s). Accuracy rose to 81.56%, nominally +1.06 over continued SFT. But McNemar gives $p = 0.678$ and the interval is $[-2.13, +4.26]$: this is indistinguishable from the registered -1.06 , and $P(\Delta \geq +3) = 0.12$. The sign moved; nothing else did.

More compute did not make the RL arm more efficient. The zero-advantage fraction *rose* with training, from 73.7% to 86.0%, as the policy solved more of the pool outright. Informative prompts grew only $5.6\times$ across a $10.6\times$ increase in steps. The win at ten times the budget was bought by brute force, not by better use of the signal.

G.3 What this does to the paper

It bounds the 7B claim and leaves the rest standing.

The registered finding is unchanged as registered. Under the pre-registered protocol—equal wall clock for every arm—GRPO does not clear +3 at 7B, and the compute-matched follow-up confirms that this is not an artifact of the degeneracy we diagnosed. Both the original -1.06 and the filtered $+1.06$ sit inside the same interval.

The 7B configuration no longer supports a claim about scale. It was included to answer “your models are too small,” and it cannot: at $10\times$ the budget the same cell reverses to a decisive Go. What Config 3 demonstrates is a *budget* boundary, not a capability one, and Section 6 now says so. The quantitative core of this paper is the 11 groups at 0.5B and 1B, where $\max \mathcal{B}$ is attained by a deterministic arm leading by 8.5 to 36.9 points rather than by 1 to 4, and where GRPO had 0.53 to 0.79 epochs rather than 0.13.

The paper’s own thesis absorbs this. The claim is that a reported RLVR effect depends on the comparison it is measured against. We can now put three numbers on that dependence, from the same runs and the same test set: changing which control arm is nominated moves the reported effect by 19.5 points (Section 8); changing the compute budget by an order of magnitude moves it by 7.8; and changing the inference engine alone moves it by 3.0 (Appendix F.3). A reader who takes only one thing from this paper should take that an RLVR number without its comparison, its budget and its decoding stack attached is not yet a result.




Figure 1: The five arms and the decision rule. Every arm branches from the same warm start W_m and sees the same m demonstrations, so the comparison is about what to do next with a fixed budget. Δ takes the *maximum* over the four-arm control group B .




Figure 2: Frozen-test Δ for all 17 groups, grouped by backbone. Every value is negative and none approaches the pre-registered +3-point Go threshold (dashed). The arm attaining max $\mathcal{B}$ is annotated per band: the solver at 0.5B and 1B, continued SFT at 7B.




Figure 3: Dose-response curves (dev; three seeds at the operating budgets, one seed at the diagnostic-only budgets). Learning arms scale with the demonstration budget; the deterministic solver does not.